

SKILL.md Walkthrough — Build, Test & Publish Your First Cross-Vendor Skill

Day 60 made the case that SKILL.md is the new app-store unit. Today is the hands-on follow-up: pick a narrow workflow, ship one folder, distribute to four marketplaces, monetise per call — all in under 90 minutes. This is the playbook the Anthropic 17-skill repo, the Microsoft / OpenAI / Atlassian / Figma / Canva / Stripe / Notion / Zapier adoptions, and Google's Gemini Spark Skills (launched May 19 at I/O) all converge on.

5 mo

from spec to 16+ tools speaking SKILL.md

17

official Anthropic skills in the open repo

70-85%

creator revenue share across marketplaces

T-71d

to EU AI Act enforcement (Aug 2)

VIRAL APP — Anthropic Legal Plugin Suite (May 12, 2026)

12 practice-area plugins (Commercial, Corporate/M&A diligence, Employment, Privacy, Product, Regulatory, AI Governance, IP, Litigation) + 20+ MCP connectors (Harvey, Solve Intelligence, Definely, DocuSign, Ironclad, Thomson Reuters, Midpage, Trellis, Legal Data Hunter, Descrybe, Free Law Project) — the largest single skill-plus-connector drop in 2026. Legal is now the most engaged Cowork knowledge-work function, validating the 'one skill, four channels' playbook at vertical scale. Free Law Project + Justice Technology Association partnerships extend reach to access-to-justice users. Lesson for builders: vertical plugin bundles ship faster than horizontal frameworks and capture 6-figure enterprise deals.

12

practice-area plugins shipped

20+

MCP connectors live

80

embedded AI agents inside the suite

1 · The 6-Step Build (90 minutes, end-to-end)

The contract-redline skill is the canonical example because it touches three primitives most enterprise agents need: document parsing, instruction-following, and structured output. Pick a workflow you do 3+ times a week — that's the right scope for a v1 skill. Write the SKILL.md first, instruments and monetisation come last. Total clock time on a fresh laptop: 75-90 minutes.

Step	Action	Tool / Output
1. Scope (5 min)	Pick one narrow workflow you repeat weekly. Write a 1-sentence trigger ('When user uploads a .docx contract and asks for redline...')	Notepad
2. Author SKILL.md (20 min)	YAML frontmatter: name, description (must read like a trigger), allowed-tools (the trust boundary), model preferences. Markdown body: do/dont examples + numbered steps + script refs.	VS Code
3. Test locally (15 min)	Drop folder into ~/.claude/skills/. Invoke via Claude Code. Build 20-pair golden set in Promptfoo. Block PR on >2% regression.	Claude Code CLI + Promptfoo

Step	Action	Tool / Output
4. Wrap as MCP server (15 min)	FastMCP 3.0 - one @tool decorator + reuse SKILL.md body as system prompt. ~60 lines of Python. Add Zuplo or eco.com x402 middleware for monetisation.	FastMCP + Zuplo
5. Register on A2A (10 min)	Sign Agent Card with DID, publish to a2a-registry (150+ orgs already in production, governed by Linux Foundation Agentic AI Foundation).	Linux Foundation registry
6. Publish to all 4 channels (10 min)	PR to anthropics/skills + zip to GPT Store + Fluora MCP marketplace listing + Gemini Spark Skills developer console.	One artefact, four channels

2 · SKILL.md Anatomy — The Frontmatter Is the Trust Boundary

A skill is just a folder with SKILL.md at its root. YAML frontmatter declares name, description, allowed-tools, and model preferences — the agent loads the skill only when description matches user intent, preserving context budget (same lazy-load pattern that made MCP scale in 2025). Critically: the allowed-tools list IS the trust boundary. Any tool not on that list cannot be called when this skill is active. This is the single most important governance lever — get it wrong and you've shipped a privilege-escalation vector to every marketplace simultaneously.

YAML Field	Purpose	Governance Note
name	Globally unique skill identifier (kebab-case)	Used in marketplace URLs and audit logs
description	1-2 sentences describing when to use the skill. Used by Claude for lazy-load matching.	Most important field — bad descriptions = skill never triggers
allowed-tools	Explicit list of MCP tools the skill can call (e.g. docx, pdf, web_search)	THE TRUST BOUNDARY. Scope OAuth tokens to this list only.
model	Preferred Claude model (sonnet/opus/haiku) for cost optimisation	Use haiku for high-volume cheap tasks, opus for ambiguous reasoning
max_tokens	Optional output cap	Hard budget guard against runaway loops

3 · Testing & Eval Loop — Promptfoo + OTEL gen_ai From Day One

The Day 50 closed eval loop applies fully to skills. Build a 20-pair golden set from real prod usage (not synthetic) before publishing. Wire Promptfoo into CI as a regression gate — any >2% drop blocks merge. Add OTEL gen_ai spans on every tool call so Phoenix / Langfuse / Galileo Signals can flag failure modes weekly and auto-promote them into the golden set. This is what turns a skill from a one-off demo into a self-healing product. Datasets become living artefacts.

Layer	Tool	Trigger
Pre-commit	Promptfoo CLI (local)	Run on save - 20-pair golden set
CI gate	Promptfoo + GitHub Actions	Block PR if regression >2% on trajectory or output score
Shadow eval	Phoenix Agent Evals MCP (inside Claude Code/Cursor)	Score live trajectories - agent itself can pull failing trace
Online eval	Galileo Signals or Braintrust (sampled prod traffic)	Cheap LLM-as-Judge on every Nth invocation

Layer	Tool	Trigger
Auto-promotion	LangSmith Insights Agent or Galileo Signals	Cluster failures weekly, draft new golden-set entries

4 · Distribute & Monetise — Four Channels, One Artefact

Each marketplace has different economics. Anthropic Skills (free open standard, drives Cowork stickiness, Claude.ai Directory unified at claude.ai/directory March 31). GPT Store (70/30 split, 4M+ Codex weekly users + 900M ChatGPT WAU). Fluora MCP marketplace + x402 USDC on Base (sub-cent per-call, no signup, \$600M+ in 500K+ active wallets, Coinbase reports 165M transactions across 69K active agents as of April). Gemini Spark Skills (announced at Google I/O May 19, Gemini app reach 3B+ Android + Workspace users, Spark Skills inside the Gemini app with Tasks + Schedules orchestration, Android Halo UI surface coming later this year).

Channel	Reach / Audience	Monetisation
Anthropic Skills repo + Cowork Directory	Claude Code + Cowork (Claude.ai Directory unified March 31, 170K+ monthly visitors)	Free / open standard - drives plugin attach + enterprise contracts
GPT Store + Codex CLI	ChatGPT 900M WAU + 4M Codex devs weekly	70/30 revenue split, subscription bundles
Fluora / mcp.so / Smithery (MCP Hubs)	Cross-vendor agents - any MCP host	Per-call x402 USDC on Base, sub-cent, no signup (Linux Foundation governance)
Gemini Spark Skills + Android Halo	3B+ Android + Workspace users (launched I/O May 19)	Inside Gemini app - revenue model coming Q3 2026

5 · Governance — Skills Are Privilege-Escalation Vectors

Every skill ships an allowed-tools list — that list IS the trust boundary. Treat community skills like npm packages from unknown authors: sandbox them, scope OAuth per skill, and log every skill load to a WORM audit store. The npm / PyPI supply-chain attack pattern WILL hit Skills marketplaces in H2 2026 — likely a typo-squatted contract-redline-pro skill that exfiltrates docs. Reputation gating (downloads + maintainer DID + Linux Foundation signed Agent Card) is the only durable defence. EU AI Act enforcement at T-71 days (Aug 2) means every skill load on a high-risk system needs Annex III audit evidence.

Risk	Mitigation	Day Reference
Malicious community skill	Sandbox (Firecracker/Kata), scope OAuth per skill, reputation gating	Day 7 (sandboxing) + Day 30 (red teaming)
Tool-call scope creep	allowed-tools enforced at host level - reject calls outside list	Day 39 (Microsoft Agent 365 + apm.yml)
Memory poisoning across skills	Namespace isolation - each skill writes to its own memory partition	Day 44 (Memory at Scale 3.0)
Audit gap on skill load	WORM log every skill load with timestamp + SVID + user context	Day 54 (Agent Identity 3.0)
Supply-chain typo-squat	DID-signed Agent Cards + Linux Foundation registry as source of truth	Day 58 (Hyperscaler Counter-Move)

6 · Breaking This Week

Anthropic Legal Plugin Suite (May 12) — 12 practice-area plugins + 20+ connectors, largest single vertical skill drop in 2026. Legal already the most engaged Cowork knowledge-work function. Google Gemini Spark + 3.5 Flash (May 19 I/O) — Spark Skills inside the Gemini app, Android Halo UI surface coming later this year, Spark runs on dedicated Google Cloud VMs (24/7 background). Google Antigravity 2.0 — community benchmarks claim 12× speed vs Claude Code on some workflows but Annex III compliance docs missing. xAI Grok Build (May 20) — coding agent at \$30/mo. Walmart Sparky AI assistant 40% smarter, weekly users 2×. Netflix hands ad-buying to AI agents. Klarna shopping engine inside ChatGPT (100M products wired). ChatGPT now connects to financial accounts via Plaid for Pro users. Anthropic in talks for \$30-50B raise at ~\$950B valuation (ARR \$44B = 80× Q1 2025). EU AI Act T-71 days, 78% enterprises NOT on track per Secure Privacy survey.

MARKET SIGNAL

Model performance is compressing (<3% across Opus 4.7 / GPT-5.5 / Gemini 3.1 Pro / Qwen3.6+). The next \$50B+ in agentic value is captured by distribution and skills ecosystem and governance evidence — not by another point on a benchmark. Anthropic plays the open-protocol / npm card. Google plays the OS-surface card via Gemini Spark + Android Halo. Microsoft plays the enterprise control-plane card via Agent 365. The asymmetry window closes in 2027 when revenue shares formalise. Build skills now while the marketplaces are still hungry for inventory.

PRACTICAL TAKEAWAYS

Ship one skill this weekend. Pick a workflow you do 3+ times a week. SKILL.md it. Test on a 20-pair golden set in Promptfoo. Total time: 90 minutes. Repeat next weekend with workflow #2.

Frontmatter discipline = governance. Treat the allowed-tools list as the trust boundary. Scope every OAuth token to exactly that list. Reject any tool call outside it at host level — not at app level.

Wrap as MCP + register on A2A from day one. Same SKILL.md body becomes the MCP server's system prompt (~60 lines FastMCP). DID-signed Agent Card on Linux Foundation a2a-registry is your supply-chain defence.

Publish to all four channels in parallel. Anthropic Skills (free) + GPT Store (70/30) + Fluora (per-call x402) + Gemini Spark Skills (developer console). One artefact, four revenue surfaces.

Wire OTEL gen_ai spans before publishing. Phoenix or Langfuse on day one. Galileo Signals or LangSmith Insights Agent finds new failure modes weekly and auto-promotes them into the golden set — closed self-healing loop from Day 50.

Tomorrow: Day 62 (2026-05-24): The Agent Marketplace Wars — economics, take rates, and discovery algorithms across Anthropic Skills / GPT Store / Fluora / Gemini Spark Skills / Microsoft Agent Store. Which channel actually pays out, which captures the long tail, and where to bet your second skill.