DAY 63 | MAY 25, 2026

The Agentic AI Valuation Framework

How to value an agentic AI company when the model is no longer the moat — DCF, comps, and the new rules driving Anthropic toward a $950B raise.

$950B	$44B	10-50x	99.9%
Anthropic talks valuation	Anthropic ARR (80x Q1 '25)	AI startup ARR multiple	Reliability that unlocks 20x

1 · Why valuation suddenly matters

Anthropic is in talks to raise $30-50B at a pre-money valuation of $900-950B — pushing it past OpenAI's $850B and within sight of a $1T mark OpenAI itself just targeted in its May 22 S-1 filing. Anthropic's ARR has compressed from $9B (late 2025) to $44B (May 2026), an 80x lift in five quarters. Cowork, Claude Code, and the 10-agent Finance pack have collapsed the per-seat SaaS model into per-task economics — and that breaks every valuation framework built around DAU and ARPU.

> ***So what:*** *Every CFO, investor, and founder is being asked the same question this quarter — what is an agentic AI company actually worth, and how do you defend the multiple when the underlying model is interchangeable?*

2 · DCF rebuilt for per-task economics

The classic SaaS DCF leans on seats × ARPU × net retention. Agentic businesses don't sell seats — they sell completed tasks. Rebuild the model around four primitives:

Primitive	Definition	2026 benchmark
Active Agents (AA)	Distinct agent identities (SVID/DID) calling your platform per month.	Anthropic: ~1.5M Claude Code dev agents, Cowork undisclosed
Tasks per Agent (TPA)	Completed multi-step tasks per agent per month. 5-50x richer than chat turns.	Salesforce Agentforce: 18,500 customers, ~190 tasks/agent/mo
Revenue per Task (RPT)	Blended price across model routing tiers (nano/mid/frontier).	$0.05-$0.50 typical; legal redline $2.00; payroll recon $0.30
Task Completion Rate (TCR)	% of attempted tasks completed without human escalation. The 99.9% threshold.	Best-in-class: 98-99.5%. The 9.9% gap is where 90% of valuation lives.

Forward revenue then becomes **AA × TPA × RPT × TCR**. Discount it at 12-15% (frontier model risk + EU AI Act enforcement risk + agent-economy regulatory tail). The terminal multiple is no longer 8-10x SaaS FCF; it's anchored to the marketplace position you've earned in the skills ecosystem (Day 60-62) — typically 15-25x for skill-ecosystem leaders.

3 · The 2026 comp sheet

ARR multiples have widened — investors have stopped paying 50x just for the AI label. The new ceiling is reserved for companies that own a workflow *and* the data feedback loop that makes the agent uniquely reliable.

Company	ARR (run-rate)	Implied valuation	ARR multiple
Anthropic	$44B	$950B (talks)	~22x

Company	ARR (run-rate)	Implied valuation	ARR multiple
OpenAI	$25B	$850B-$1T (S-1)	~34-40x
xAI	~$5B	$250B (SpaceX deal)	~50x
Cursor	$2B	$50B (rumoured)	~25x
Databricks	$4.8B	$134B	~27.9x
DeepSeek	Undisclosed	China-only ~$80B	n/a
Cohere	$200M	$5B	~25x
Mistral	$120M	$11.5B	~96x (sovereignty premium)

Two patterns are visible. First, frontier-lab valuations have decoupled from current ARR and re-anchored to **distribution surface** (OpenAI's 900M WAU, Anthropic's enterprise share). Second, EU-sovereign players (Mistral) trade at extreme premiums on geography alone — a direct consequence of the August 2 EU AI Act enforcement deadline now T-69 days away.

4 · The new moat: skills + governance, not model

Model performance has compressed to under 3% across Claude Opus 4.7, GPT-5.5, Gemini 3.5 Flash, and Qwen3.7-Max. That means the model is no longer a moat — it's a commodity component inside model-routing stacks (Day 43). The valuation multiple now lives in three places:

Moat layer	What it looks like	Multiple impact
Skills ecosystem	Number of installed SKILL.md packages, A2A registry signed agent cards, MCP server installs. Anthropic Skills repo + 16 tools speaking the spec.	+5-10x on base ARR multiple
Governance evidence	Pre-built Annex III audit trails, SPIFFE/SVID identity, T1-T4 kill switches, NIST CAISI WORM logs. Microsoft Agent 365 is the reference.	+3-7x for regulated verticals
Distribution lock	Native placement inside an OS surface (Gemini Spark, Apple Intelligence) or enterprise control plane (Agent 365, Salesforce Agentforce).	+5-10x driven by switching cost
Reliability proof	TCR > 99% sustained across millions of tasks. The 9.9% gap between 'works most of the time' and 'works always' is where 90% of valuation lives.	Cap at 3-4x if reliability cannot be proven

5 · Treating agents as customers with LTV

The biggest break with SaaS valuation is that **agents are now customers**. Coinbase reports 165M x402 transactions across 69K active agent wallets in April 2026; $600M+ now sits in 500K+ active agent wallets. These agents pay, renew, churn, and refer other agents — and they need their own LTV/CAC model:

Metric	Human SaaS analogue	Agent-economy translation
LTV	ARPU × gross margin / churn	RPT × tasks/month × expected agent lifetime (rolling 30d quality score)
CAC	Sales & marketing / new logos	Marketplace placement fees (Fluora/mcp.so leaderboard cost) + SDK adoption cost

Metric	Human SaaS analogue	Agent-economy translation
Net retention	Expansion / churn	Tool-call expansion as orchestrator agents discover more of your skills
Churn signal	Login drop-off, NPS	OTEL gen_ai span drop, cosine<0.7 drift, model rotation by orchestrator

> ***Practical takeaway:*** *If your forecast still slices revenue by 'paid seats', you're under-valuing every multi-agent customer by 5-20x. Re-slice by tasks billed and agent identities served — the same artefact that produces your EU AI Act Annex III audit pack.*

6 · Pure-skill, pure-marketplace, pure-orchestration

Three new startup archetypes need bespoke valuation frames:

Archetype	Revenue model	Multiple range	Closest comp
Pure-skill	Per-call x402 USDC on Base; 70-85% creator share via marketplaces.	8-15x net revenue	Hermes Agent skill libraries; legal-redline templates
Pure-marketplace	2.5% transaction fee (volume-discounted to 1.5%) on agent payments.	20-35x net revenue (x402 leaderboard premium)	Fluora, mcp.so, Smithery
Pure-orchestration	Per-task fee on multi-agent workflows; LangGraph-style platform.	12-22x ARR	CopilotKit, LangChain (1B+ downloads)

7 · Public market read-through

Three public tape signals validate the framework: (a) Microsoft's M365 E7 tier — the first new enterprise SKU since E5 (2015) — added a measurable ARPU lift credited explicitly to Agent 365 in MSFT's Q3 FY26 call. (b) Salesforce traded to a fresh agentic premium after Agentforce 3.0 crossed $500M ARR with 18,500+ customers. (c) ServiceNow's Autonomous Workforce drove gross margin expansion of ~180 bps QoQ as 90%+ of IT tickets moved off human agents. Public markets are now paying explicitly for the same three moats — skills ecosystem, governance evidence, and distribution lock — that private agentic comps trade on.

8 · Viral AI app of the day

OpenHuman by tinyhumansai — climbed to #1 on GitHub Trending May 13-16 by inverting the agent UX playbook. Instead of waiting for a first prompt, it pre-loads a 1B-token Memory Tree from 118+ services via one-click OAuth (Gmail, GitHub, Slack, Calendar, Drive, Stripe, Linear, Jira). A 'subconscious loop' runs every 20 minutes without user input, reading to-dos and recent memory to decide what to act on. The desktop mascot can join Google Meets as a separate participant, transcribing into the Memory Tree and speaking back. All local-first (Rust + Tauri, MIT, SQLite + Obsidian wiki) — no data egress unless explicitly escalated.

> ***Valuation read:*** *OpenHuman is a pure-skill / pure-orchestration hybrid with zero ARR today — but the 1B-token Memory Tree is exactly the kind of data feedback loop investors are paying 15-20x ARR for. Watch for a paid SaaS tier or A2A registry play within 90 days.*

9 · Practical takeaway for tomorrow

If you're a CFO, build a one-page agentic valuation deck with five blocks: (1) AA × TPA × RPT × TCR forward revenue, (2) ARR multiple anchored to 22x (Anthropic comp) discounted by your skills ecosystem position, (3) Skills + governance + distribution moat scorecard, (4) Agent-LTV/CAC for any third-party agent integrations, (5) Annex III governance evidence as a discount-rate reducer. If you're a founder, publish a SKILL.md into all four marketplaces this week — the asymmetry window closes in 2027 when revenue shares formalise.

Tomorrow (Day 64): *Agentic AI in Financial Services — Anthropic's 10 finance agents in production, JP Morgan's COiN at scale, and the new FinCEN AML rules landing Q3 2026.*