

The EU AI Act August 2 Enforcement Playbook

H2 opens with the regulatory event the whole series has been building toward. **August 2, 2026** is when the EU AI Act's next wave of obligations becomes enforceable — and this year the date comes with a twist. A last-minute **Digital Omnibus** deal (provisional May 7) would push the heaviest **high-risk (Annex III)** obligations out to **December 2, 2027**, but it is **not yet law** and several obligations — chiefly **Article 50 transparency**, the governance machinery, and GPAI penalties — stay live on August 2 regardless. This playbook explains exactly what is binding on the day, what the high-risk evidence pack looks like once it lands, the penalty tiers, and how to turn an OTEL + WORM audit trail into a clean regulator artefact. The throughline of the whole series holds: *a governed agent you can audit beats a smarter agent you can't* — and the audit is now a legal deliverable, not a nice-to-have.

Aug 2 2026 — Article 50 transparency + governance + GPAI penalties go live	Dec 2027 where high-risk (Annex III) moves IF the Omnibus is adopted — not yet law	EUR 35M / 7% top penalty tier (prohibited practices) — whichever is higher of fine or global turnover	78% of enterprises not meaningfully prepared for the obligations (2026 readiness survey)
--	--	--	---

1 · Read the date correctly — August 2 is now a split deadline

For two years “August 2, 2026” has been shorthand for “the day high-risk AI gets regulated.” That is no longer the clean story. On **May 7, 2026** the Council, Parliament and Commission reached a **provisional political agreement** on the Digital Omnibus, which **defers stand-alone Annex III high-risk obligations to December 2, 2027** (and AI embedded in regulated products under Annex I to August 2, 2028). But that deal only takes legal effect once it clears the final European Parliament plenary vote (expected in June), is formally adopted by the Council, and is **published in the Official Journal** — expected, but not guaranteed, before August 2. **Until publication, the original August 2, 2026 date formally stands.** So you must hold two clocks at once: what is binding on the day no matter what, and what is ‘probably’ deferred but could snap back if the Omnibus slips.

Obligation	Status on Aug 2, 2026	Why it matters
Article 50 transparency (chatbot disclosure, AI-content marking, deepfake labelling)	LIVE — not deferred	Any agent that talks to people or generates synthetic media must disclose + (machine-readable) mark its output. This is the obligation most teams actually touch.
Governance machinery (AI Office, national competent authorities, penalty regimes in each Member State)	LIVE	Regulators get teeth on the day: market surveillance, complaint handling, the power to demand documentation.
GPAI penalties (Article 101, enforced by the AI Office)	LIVE	Foundation-model providers face EU-level fines from Aug 2; pre-Aug-2025 models have until Aug 2027 to comply.
High-risk Annex III obligations (risk mgmt, technical docs, logging, human oversight, conformity assessment, EU-DB registration)	Proposed move to Dec 2, 2027 — but ONLY once the Omnibus is published; otherwise live Aug 2	The heavy evidence pack. Treat the relief as provisional: keep building the pack on the original clock until the OJ publication is real.

Obligation	Status on Aug 2, 2026	Why it matters
New Article 5 prohibitions (non-consensual intimate imagery / 'nudifiers', CSAM generation)	Added by the Omnibus — stricter, not looser	The deal simplifies timelines but tightens the red lines — 'delay' is not 'deregulation'.

***The one mental model:** August 2 didn't get cancelled — it got split. Transparency, governance and GPAI penalties are a hard deadline; the high-risk evidence pack is 'probably 16 months later, but build it now anyway.' Plan to the earlier of the two clocks and you are never caught out.*

2 · Article 50 — the obligation almost every agent touches Aug 2

Because the heavy high-risk pack is (probably) deferred, **Article 50 transparency is the live obligation most teams will meet first**. It is a two-sided duty. **Providers** of generative systems must ensure outputs are **machine-readable as AI-generated** — watermarking plus C2PA / IPTC metadata, the same provenance stack covered on Day 79 (media) and Day 34 (multimodal). **Deployers** must add **visible disclosure**: tell people they're talking to an AI at first interaction, and clearly label deepfake image / audio / video as artificially generated, at the latest on first exposure. Narrow carve-outs exist for law-enforcement and "obviously artistic or fictional" work. A draft Code of Practice gives the practical recipe: imperceptible technical marking from the provider, plus a visible "AI" label in a fixed position from the deployer.

If your agent...	Article 50 duty on Aug 2	How to satisfy it
Talks to people (chat, voice, email outreach, support)	Disclose it's an AI at first interaction	Persistent "AI" badge / opening disclosure wired into the run; log the disclosure event (ties to the AG-UI RUN_STARTED pattern, Day 48).
Generates images / audio / video / synthetic voice	Machine-readable marking (provider) + visible label (deployer)	SynthID-class watermark + signed C2PA manifest on every asset; visible "AI-generated" label. Provenance = the trust boundary.
Writes text published to the public (articles, posts, ads)	Disclose AI-generation where it informs the public	Label AI-authored public-facing text; keep the prompt-to-output trail as evidence.
Is purely internal / analytical (no public output, no human-facing persona)	Largely out of Article 50 scope	Still log it — you'll want the trail for the high-risk pack and for governance requests.

3 · The high-risk evidence pack — what Annex III actually demands

If your system lands in **Annex III** (biometrics, critical infrastructure, education, employment / HR, essential public & private services incl. credit and insurance, law enforcement, migration, justice) you owe a full conformity story. The Omnibus likely buys you until **Dec 2027**, but the pack takes 12–18 months to build — and 61% of enterprises have **no process** to produce it — so the start date is now, not the deadline. The deliverable is the **Annex IV technical file**: nine sections demonstrating compliance, kept *live* (a maintained record, not a write-once document) and handed to national authorities on request. Most Annex III systems can **self-assess**; only remote biometric identification and certain law-enforcement uses need an external **notified body**. After assessment you issue an EU Declaration of Conformity, apply CE marking, and **register the system in the EU database** before it goes live.

Annex IV / Article	What the regulator wants to see	Agentic-stack source (where it already lives)
System description & intended purpose	What the agent does, its boundaries, foreseeable misuse	apm.yml manifest + SKILL.md allowed-tools list = the declared scope (Day 39, Day 60).

Annex IV / Article	What the regulator wants to see	Agentic-stack source (where it already lives)
Risk management (Art. 9)	Risk register: probability × severity, all health / safety / rights / discrimination risks, mitigation traceability, residual-risk sign-off	Threat model + kill-switch design + red-team results (Day 21, Day 54).
Data governance	Training / input data provenance, quality, bias controls	Memory IDs + namespace isolation + data-residency story (Day 44, Day 55).
Logging / record-keeping (Art. 12)	Automatic, lifetime event logging that reconstructs each decision	OTEL gen_ai spans → WORM store — the evidence pipeline (Day 22, Day 50).
Human oversight (Art. 14)	How a human can understand, override, and stop the system	AG-UI approval gates + T1–T4 kill switch <1s (Day 48, Day 49).
Accuracy, robustness, cybersecurity + Declaration of Conformity + post-market monitoring	Performance metrics, security posture, signed DoC, ongoing monitoring plan	Eval golden sets + online evals + SPIFFE/SVID identity + drift alerts (Day 45, Day 54).

The reframe that makes this cheap: you are not building compliance from scratch — you are relabelling artefacts you already produce. The OTEL trace is the Article 12 log. The kill switch is the Article 14 oversight control. The apm.yml is the system description. Build the agent right and the evidence pack falls out as a by-product.

4 · Penalties, and the audit pipeline that doubles as your defence

The Act's fines are deliberately existential, and use a “whichever is higher” rule (fixed sum or percentage of global turnover). **Prohibited practices: up to EUR 35M or 7%. High-risk & GPAI breaches: up to EUR 15M or 3%. Misleading authorities / false documentation: up to EUR 7.5M or 1%.** Enforcement splits between the **AI Office** (general-purpose models) and **national competent authorities** (everything else) — and both go live on August 2. The single most leveraged thing you can do before the date is wire one pipeline: **OTEL gen_ai spans** → **WORM (write-once) audit store**. Every model call, tool call, memory write, approval gate and kill-switch event lands as an immutable, timestamped record tagged with the agent's SPIFFE identity. That same trace is your Article 12 log, your Annex III evidence, and — if a regulator ever asks “why did the agent do that?” — your reconstructable answer.

Violation tier	Maximum penalty	Who enforces
Prohibited AI practices (Article 5 red lines)	EUR 35M or 7% of global annual turnover	National competent authorities
High-risk breaches + GPAI obligations	EUR 15M or 3% of global turnover	National authorities / AI Office (Art. 101 for GPAI)
Supplying false / misleading info to authorities	EUR 7.5M or 1% of global turnover	National competent authorities

***90-day sprint to Aug 2: (1) Inventory** every AI system + classify it (prohibited / high-risk / limited / minimal) — 50%+ of firms can't yet, and you can't comply with what you can't see. (2) **Ship Article 50** — disclosure + watermarking on every public-facing agent (this is the live obligation). (3) **Stand up the OTEL→WORM pipeline** and start the Annex IV file on the original clock — the same trail is evidence, log and liability firewall in one.*

5 · Viral app spotlight — OpenCode, the most-starred open agent

This week's standout is **OpenCode** (opencode.ai, by the SST team at Anomaly Innovations), which crossed **~172,000 GitHub stars** to become the most-starred open-source AI agent — ahead of Gemini CLI and Codex — with ~7.5M monthly active developers, ~900 contributors and 13,000+ commits in under a year. Its thesis is the perfect coda to a governance issue: **the model is a pluggable dependency, not the product**. OpenCode is a pure terminal agent (no IDE extension, no web UI) that runs on Claude, GPT-5, Gemini, Grok, Bedrock, Azure OpenAI or a local Ollama model — and it **stores none of your code or context**, which is exactly the data-residency posture EU deployers need. It is the clearest proof of the Day 80 compression thesis: when frontier models sit within ~3% of each other, the durable value moves to the layer you control — the harness, the data boundary, the audit trail. (OpenClaw still tops the raw OSS charts at 210K+ stars, now governed by a foundation.)

~172K GitHub stars — most-starred open-source AI agent in 2026	7.5M monthly active developers in under a year	0 code / context bytes stored — model-agnostic, privacy-first by design
--	--	---

Practical takeaways

1. Hold two clocks — and plan to the earlier one.

Article 50 transparency, governance and GPAI penalties are **binding on August 2, 2026**. High-risk (Annex III) is *probably* deferred to Dec 2027 by the Digital Omnibus — but only once it is published in the Official Journal, which hasn't happened yet. Build the high-risk pack on the original clock so a slipped Omnibus never catches you out.

2. Ship the live obligation first: disclose + watermark.

Every public-facing agent needs AI disclosure at first interaction and machine-readable marking (SynthID-class watermark + signed C2PA manifest) on synthetic media — now. This is the obligation most teams actually touch on August 2, and it's a days-not-months fix.

3. Wire OTEL → WORM once; reuse it as evidence, log and defence.

Stand up a single pipeline that lands every model call, tool call, approval gate and kill-switch event as an immutable, SPIFFE-tagged record. That same trail is your Article 12 log, your Annex IV evidence, and your reconstructable answer to “why did the agent do that?” Build the agent right and compliance falls out as a by-product.

***Tomorrow (Day 82):** The Agentic AI Procurement Playbook — how governance becomes a contract line item. Buyer checklists for SPIFFE/SVID identity, kill switches, C2PA/SynthID, Annex IV evidence and SLAs; the build-vs-buy decision for the control plane; and why ‘show me your audit trail’ is the new RFP question.*