

The Super-App Wars: ChatGPT Work vs. Claude Cowork

July 9 became the first day three frontier labs shipped at once — and the real fight moved from chatbot to always-on work agent

On July 9, something happened that had never happened before: OpenAI, xAI, and Anthropic all pushed major frontier updates within the same 48 hours. OpenAI merged ChatGPT and Codex into one desktop app and launched a new “Work” agent tier powered by GPT-5.6 (Sol / Terra / Luna); xAI shipped Grok 4.5 into the wild; and Anthropic took Claude Cowork out of desktop-only beta onto web and mobile. The benchmark leaderboards are the story everyone is watching — Grok 4.5 landed 4th on Artificial Analysis's Intelligence Index, behind Claude Fable 5, GPT-5.5, and Opus 4.8 — but the real story is what all three labs agreed on without coordinating: the product surface is moving from a chat window you type into, to a work agent that runs for hours while you do something else, and you only check in when it needs you. 104 days into this series we've covered the frontier-model race compressing to a near-tie (Day 80), self-improving agents (Day 100–102), and 24+ industries running agents in production — today is about the surface war being fought on top of all of it.

53.6 GPT-5.6 Sol's new high on Agents' Last Exam, the agentic-reasoning benchmark	#4 Grok 4.5's rank on Artificial Analysis's Intelligence Index, behind Fable 5, GPT-5.5, Opus 4.8	1.2M Cowork sessions Anthropic analyzed across 600,000+ organizations	33.4% share of that usage that was plain business-process work, not coding
---	---	---	--

1) Three labs, one bet: the chat window is becoming a work agent

GPT-5.6 ships in three tiers that are worth understanding on their own: Sol is the flagship, tuned for the hardest multi-step agentic reasoning; Terra matches the older GPT-5.5 at roughly half the cost, for everyday production use; Luna is the fastest and cheapest, for high-volume routine calls. That's the same match-the-tier-to-the-job logic this series has covered before — the news is that OpenAI wrapped all three inside one merged desktop app, alongside Codex (coding) and a new Work tab, and made Chat, Work, and Codex free on every plan, including Free. Anthropic's move is different in shape: Claude Cowork, which launched desktop-only in January, is now on web and mobile too — but deliberately as a monitor, not a remote control. You can start a task at your desk, get a phone notification when it needs a decision, and pick up the finished output later; the agent still executes on your own machine, not on the phone itself.

Lab	What shipped July 9	Where the agent executes	Free to use?
OpenAI	ChatGPT + Codex merged into one app; new Work tab; GPT-5.6 (Sol/Terra/Luna)	In-app, cloud-side	Yes – every plan, incl. Free
Anthropic	Claude Cowork expands from desktop-only to web + mobile	Your own machine; phone only monitors	No – Max subscribers, beta
xAI	Grok 4.5 general release	Chat / API	Yes, with usage limits

2) What 1.2 million real sessions reveal: almost nobody is using this for code

Anthropic published a breakdown of 1.2 million real Cowork sessions across more than 600,000 organizations, and it cuts against the industry's own marketing. Software development — the use case behind almost every coding-agent headline — was just 8.7% of sessions. The largest category, at 33.4%, was plain business-process work: pulling scattered updates into a status report, reconciling a spreadsheet, chasing down a number across five tabs. Content creation and copywriting came in second at 16.4%. The lesson isn't that coding agents don't matter — it's that the first place a general-purpose work agent earns its keep inside a real organization is the boring, recurring admin work nobody wanted to automate first, not the flashiest demo.

Session category	Share of 1.2M sessions	What it looks like in practice
Business process work	33.4%	Reconciling spreadsheets, assembling a status report from scattered updates
Content creation / copywriting	16.4%	Drafting marketing copy, decks, and long-form writing
Software development	8.7%	Writing, reviewing, or shipping code

3) Which work-agent pattern to reach for

A decision guide. Reach for a merged chat-plus-work-plus-code surface (ChatGPT Work/Codex-style) when you want one app your whole team already opens daily, and you're comfortable scoping each mode's permissions separately. Reach for a monitor-from-anywhere pattern (Cowork-style) when the agent needs to run for hours unsupervised and you want a way to check on it without giving it more places to act from. Reach for the top model tier (Sol-class) only for genuinely hard multi-step reasoning — default to the cheaper mid or entry tier for everyday and high-volume work.

When to reach for it	Use	Guardrail
Need one surface for chat, code, and long-running tasks	Merged app pattern (ChatGPT Work/Codex)	Scope each mode's permissions separately — don't let Work tasks inherit Codex's repo access
Need to delegate and step away, then check in remotely	Monitor-from-anywhere pattern (Cowork mobile)	Keep the phone surface read-only until you trust the guardrails enough to extend it
Need to solve a genuinely hard multi-step task	Top model tier (Sol-class)	Budget for the cost; don't default to it out of habit for routine work
Need cheap, high-volume routine calls	Entry tier (Luna-class)	Route by task complexity, not by whichever model you already have open

THE ONE RULE

An agent you can't supervise from outside your desk is an agent you'll eventually stop supervising. Anthropic's choice to make Cowork mobile a monitor, not a remote control, is the tell: the safest place for a long-running agent to execute is the machine you already trust, and the safest way to extend its reach is to extend your visibility into it — not its permissions.

VIRAL / OPEN-SOURCE SPOTLIGHT

ChatGPT Work goes free on every plan — including Free

Overnight, hundreds of millions of existing ChatGPT users saw a new Work tab appear next to Chat and Codex — no download, no waitlist, no separate account. Work is an agent that gathers context across your connected apps, breaks a goal into steps, and returns a finished doc, sheet, slide, or small web app, staying with a complex project for hours by completing pieces of it independently. The idea itself isn't new — Claude Cowork and Manus-style agents got there first — but shipping it by default, for free, inside the app OpenAI says has 900M+ weekly users, is a distribution move rivals can't easily match without a separate download.

Free Availability tier for Chat, Work, and Codex in the new merged desktop app	3 surfaces merged into one app: Chat, Work, and Codex	GPT-5.6 model family powering Work by default	07.09 launch date — same day as Grok 4.5 and Cowork's mobile expansion
--	---	---	--

MARKET SIGNAL

Three frontier labs releasing independently on the same day is itself the signal. With raw intelligence scores compressed to a near-tie (Day 80) and self-improvement techniques already handling incremental gains without retraining (Day 100–102), the open differentiation frontier has moved from “whose model is smartest” to “whose agent do you trust to work unsupervised, and where can you check on it.” OpenAI is betting on distribution — give Work away free inside an app hundreds of millions of people already open daily. Anthropic is betting on trust — gate Cowork’s expansion behind paying Max subscribers and a deliberately limited mobile surface. Both bets can win in different segments; the losing bet is standing still on chat-only.

PRACTICAL TAKEAWAYS

1. Separate “chat” from “work” in your own agent's UX, even inside one app

OpenAI kept Chat, Work, and Codex as separate tabs inside one merged app rather than blending them into a single mode — each has a different session length, permission scope, and failure tolerance. If you're building an agent product, resist the urge to make one mode that does everything; users need to know which mode they're in before they grant it access.

2. Design for the 33%, not just the 9%

Coding agents get the headlines, but Anthropic's own usage data shows business-process work (reconciling data, assembling reports) is roughly 4x more common than software development in real Cowork sessions. If you're prioritizing what to automate first inside your org, start with the recurring status-report-and-spreadsheet work, not the flashiest coding demo.

3. Extend an agent's reach with visibility, not permissions

When you let an agent run unsupervised for hours, the safer next step is giving yourself a way to check on it from anywhere — a phone notification, a status page — not giving the agent more places it can act from. Cowork's mobile app is a monitor, deliberately not a remote control; copy that shape before you copy the “agent that can do everything from anywhere” version.

TOMORROW

Tomorrow — Day 106: Grok 4.5 under the hood — the token-efficiency claim that changes agentic cost math, and the political-bias benchmark debate that broke out within hours of launch.