

The Good, Better, Best Playbook

OpenAI Ships Three Models at Once with GPT-5.6 — and Why Every Lab Is Copying the Split

Sol, Terra, and Luna landed in a single release on July 9, not as a future roadmap — and it's the same three-rung shape Anthropic already sells as Opus, Sonnet, and Haiku, while xAI just answered with the opposite bet: one model, priced like the cheap tier. The real story isn't a new model. It's that "one flagship" has stopped being how frontier labs compete.

Day 108 closed on a promise: OpenAI shipped three models at once with GPT-5.6 — Sol, Terra, and Luna — and today we unpack why. Meanwhile, the deadline that filled yesterday's edition arrived on schedule: China's anti-addiction rules for companion AI took effect today, and Doubao's companion agents went dark this morning, redirecting millions of users to ByteDance's Maoxiang app, our spotlight below. 109 days into this series we've watched frontier capability compress toward a near-tie (Day 80), watched three governments converge on the same companion-AI guardrails independently (Day 108), and watched Grok 4.5 undercut the market on token pricing just a day before GPT-5.6 arrived (Day 106); today the axis is pricing architecture — how many models a lab actually sells you, and what that choice is designed to make you do.

3	\$5 → \$1	2x	\$2 / \$6
Models shipped in a single GPT-5.6 release — Sol, Terra, Luna	Input price spread per million tokens, Sol (flagship) down to Luna (budget)	Terra's claimed cost efficiency vs. GPT-5.5 at comparable everyday-task performance	Grok 4.5's flat per-million-token price — one tier, undercutting rivals by 60%+

1) Three tiers, one release: what Sol, Terra, and Luna actually do

OpenAI didn't ship "GPT-5.6" as a single model — it shipped a family. **Sol** is the flagship: built for frontier reasoning and long-horizon agentic work, the tier you'd point at a hard multi-step coding task, a scientific-reasoning problem, or a security research job where getting it right matters more than getting it fast. **Terra** is the workhorse: OpenAI pitches it as matching the previous flagship's (GPT-5.5) performance at roughly half the cost, aimed at the high-volume business tasks most companies actually run all day — customer support, internal tools, document analysis. **Luna** is the cheapest and fastest: tuned for summarization, drafting, classification, and routine automation, where latency and price matter more than reasoning depth. Think of it as good/better/best, except the "best" option isn't the default anymore — Terra and Luna exist precisely so most calls never touch Sol at all.

2) The other new thing: agents that write their own coordination code

Buried inside the release is a feature that matters more than the tiering: Programmatic Tool Calling in the Responses API. Normally an agent calling tools works one step at a time — call a tool, wait, read the

result, decide the next call, repeat. Programmatic Tool Calling lets the model instead write a small in-memory program that coordinates several tool calls and processes their intermediate results itself, without round-tripping every step back through the model. In plain terms: instead of a chef shouting one instruction to the kitchen at a time and waiting for each dish before ordering the next, the chef hands over a full recipe card the kitchen can run start to finish. OpenAI is also pitching a multi-agent pattern where Sol can spin up sub-agents for parallel, focused work — the orchestrator/specialist split this series has covered before as agent orchestration frameworks matured, now built into the core API rather than something you had to hand-wire yourself.

3) Why every lab is doing this now

GPT-5.6's three-tier shape isn't new — it's OpenAI catching up to a pattern Anthropic has run for over a year with Opus, Sonnet, and Haiku, and that Google runs with its Gemini tiers. What's new is how tightly labs are now pricing each rung to force a routing decision: call the expensive model only when the cheap one demonstrably fails. xAI took the opposite bet on July 8, a day before GPT-5.6 landed — Grok 4.5 ships as one model, priced at \$2/\$6 per million tokens, positioned as "Opus-class but faster and cheaper" rather than split into rungs at all. Whether tiering or flat pricing wins depends entirely on whether most of your traffic is actually simple (tiering wins, most calls hit the cheap rung) or most of it needs frontier reasoning anyway (flat pricing wins, you're paying flagship rates either way).

THE ONE RULE

A frontier lab no longer sells you one model — it sells you a menu, and the menu is designed to make you default to the cheap end. The skill that matters now isn't picking the smartest model available; it's building the routing logic that calls the cheapest tier that clears the bar for each task, and escalates only when it doesn't.

Lab	Flagship (in / out per 1M)	Mid-tier (in / out per 1M)	Budget (in / out per 1M)
OpenAI — GPT-5.6	Sol — \$5 / \$30	Terra — \$2.50 / \$15	Luna — \$1 / \$6
Anthropic — Claude	Opus 4.8 — \$5 / \$25	Sonnet 5 — \$2 / \$10 (intro, through Aug 31)	Haiku 4.5 — \$1 / \$5
xAI — Grok 4.5	One tier, flat — \$2 / \$6 (no separate rungs)		

VIRAL / OPEN-SOURCE SPOTLIGHT

Maoxiang: the app about to absorb millions of orphaned Doubao companions

Today is the exact deadline Day 108 previewed: China's anti-addiction rules for anthropomorphic AI took effect this morning, and Doubao — a 350-million-user app — pulled its companion-agent feature on schedule. ByteDance isn't abandoning the category; it's redirecting users to Maoxiang (猫箱), a separate standalone app it already runs, built around the compliance architecture the new rules require — anti-addiction detection, an instant-exit mechanism, usage notifications — baked in from the start rather than

retrofitted. Maoxiang's monthly active users had actually been sliding, from a peak above 6 million down to roughly 3.9 million by June, before today's forced migration. It's a rare case of a regulation itself acting as the growth engine: millions of users with an emotional habit and nowhere else compliant to take it are being funneled, this week, into one specific app. Whether Maoxiang's numbers snap back above that 6-million peak by the end of July is the numbers to watch.

350M Doubao's total user base before today's companion-feature shutdown	3.9M Maoxiang's monthly active users as of June 2026, down from a 6M+ peak	1 Compliant migration path ByteDance is offering — Maoxiang, its only other companion app
---	--	---

MARKET SIGNAL

Day 108 tracked regulators converging on companion AI's business model as the risk itself. Today's signal is about a different business model — inference pricing. Three frontier releases landed inside ten days (Sonnet 5, Grok 4.5, GPT-5.6), and every one of them is a pricing move as much as a capability one: Anthropic held Sonnet 5 at intro pricing through August, xAI undercut everyone by shipping one model instead of three, and OpenAI answered by widening its own tier spread to \$1–\$30 per million tokens. When three labs reprice within the same fortnight, that's not routine — it's a live price war over exactly the workloads Terra, Sonnet 5, and Grok 4.5 are all fighting for: the high-volume, not-that-hard business tasks that make up most real usage. The flagship tiers barely moved; the mid-and-budget tiers are where the war is actually happening.

PRACTICAL TAKEAWAYS

1. Default to the middle tier, and route up only on failure

Terra, Sonnet 5, and Grok 4.5 all exist because most of what you actually call an LLM for — support replies, document summaries, internal tools — doesn't need flagship reasoning. Build your default call at the mid tier and escalate to Sol or Opus only when the cheap path demonstrably fails, rather than reaching for the smartest model out of habit.

2. If your pipeline chains more than two or three tool calls, look at programmatic tool calling

Round-tripping every tool call back through the model adds latency and tokens you're paying for on every hop. If you're hand-wiring a loop of "call tool, read result, decide next call," check whether your provider's API now lets the model write that coordination logic once instead — it's the same win hand-built multi-agent orchestration frameworks were already solving for.

3. Don't lock into one lab's tier structure — the pricing is still moving

Three repricings in ten days means whatever tier math you did last month is already stale. Keep your routing logic provider-agnostic enough that swapping which lab handles your "cheap tier" traffic is a config change, not a rebuild, because the cheapest adequate model is going to keep changing this year.

TOMORROW

Tomorrow — Day 110: the open-weight counter-punch. GLM-5.2, DeepSeek V4, and Qwen 3.6 all landed within days of the Sonnet 5 / Grok 4.5 / GPT-5.6 wave — whether free-to-run open models are actually closing the gap on the paid tiers this fast, or just closing it on paper, and what that means for which model you should be building on.

[varunsingla.com](#) · [AI Daily Learning](#) · [Day 109](#) · [The Good, Better, Best Playbook: OpenAI's GPT-5.6 Tiering, Programmatic Tool Calling & the Maoxiang Migration \(GPT-5.6 · Sol · Terra · Luna · Model Tiering · Maoxiang\)](#)