

SpaceX Tried to Buy a Second AI Coding Startup -- Its CEO Said No, Publicly

Five days after closing a \$60 billion deal for Cursor, SpaceX went shopping again -- as OpenAI's CFO tells staff to expect a public listing by 2027 and Claude starts reading raw lab instrument data for chemists

Day 145 · Thursday, August 20, 2026

Why this matters: today's three stories are all about AI companies converting momentum into permanence -- one buying its way into a market instead of building for it, one telling employees exactly when it plans to answer to public shareholders, and one proving a model can already do a piece of a trained chemist's job in under half an hour. None of these are demos anymore; they're the infrastructure decisions the next few years of AI will run on.

VIRAL SPOTLIGHT

"I Called Five Stores So You Don't Have To" -- Google's Shopping AI Starts Picking Up the Phone

Google's agentic shopping tools finished rolling out across the US this summer, and the feature people are actually talking about is the phone calls. Ask Google to find something -- a specific bike helmet, a restock of a sold-out blender -- and its AI, built on the same Duplex voice technology that once faked a human booking a haircut, will call nearby stores itself, ask about price and stock, and text you a summary of what it heard. A companion feature, agentic checkout, will also complete the purchase for you once a tracked item hits your target price, with your approval and Google Pay handling the transaction. The AI identifies itself as automated on every call, and stores can opt out of receiving them.

Why it's taking off: for a decade, "AI phone call" meant a scam robocall or an awkward customer-service bot. This is the first version millions of ordinary shoppers are hearing on the other end of a real call and coming away impressed rather than annoyed -- which is exactly why it's spreading by word of mouth instead of an ad campaign. It's also a preview of a bigger shift: the same week an AI store manager fired a human employee (covered here yesterday), the AI on the other side of that same transaction is now the one calling in.

1) SpaceX's Second Knock on Cognition's Door -- and a Public 'No'

SpaceX's \$60 billion acquisition of the AI coding startup Cursor became official on August 14. Five days later, people familiar with the matter told reporters SpaceX had also approached Cognition AI -- maker of the autonomous coding agent Devin, used inside Mercedes-Benz and GE Aerospace -- about a possible second takeover. Cognition CEO Scott Wu shut the story down within hours, posting that the report was inaccurate, that the companies "haven't been in talks," and that Cognition "is not for sale."

Why it matters: whether or not this specific approach happened, the pattern around it is real -- Elon Musk is trying to assemble AI coding capability at SpaceX by buying finished companies rather than building the equivalent in-house, the same playbook Meta used days earlier when it bought into the coding-agent market instead of shipping a homegrown one. A public, on-the-record denial from a target's CEO is also a signal in its

own right: it tells every other AI startup watching that turning down a nine-figure buyer is now a viable, even reputation-building, move.

2) OpenAI's CFO Tells Staff: Public by 2027, Maybe Sooner -- and Anthropic Could Beat Them There

In an all-hands meeting, OpenAI CFO Sarah Friar told employees the company “will be a public company in 2027,” or earlier if growth keeps “inflecting.” She framed the listing as “not a finish line, it is a milestone, another fundraise,” pointing to the \$122 billion OpenAI raised in March as proof it has room to wait. Asked about competitive pressure, Friar acknowledged Anthropic could plausibly go public as soon as September -- days after Anthropic reported its first-ever profitable quarter -- but said OpenAI is “running our own race.”

Why it matters: an IPO forces the kind of financial discipline -- clean books, predictable spending, an answer for when losses turn into profit -- that fast-growing AI labs have mostly been able to defer by raising private mega-rounds instead. Whichever of OpenAI or Anthropic lists first will be the one whose unit economics get audited in public first, and the other's playbook will likely follow whatever the market rewards or punishes.

3) Claude Just Read a Chemist's Raw Lab Files -- and Beat the Clock

Anthropic gave Claude Opus 5, working inside its new Claude Science workbench, the raw output files a contract lab produces to check whether a chemical sample is pure and correctly identified -- an NMR spectrum (which reveals a molecule's structure from how its atoms respond to a magnetic field) and an LC-MS readout (which separates and weighs the components of a mixture). Given only those files and a short plain-language prompt, Claude processed the NMR result in 23 minutes and the LC-MS result in 19, working the two in parallel, and came back within 0.1% of the purity figure the human lab reported. Anthropic says it's now preparing a broader access program so more scientists can test the workflow themselves.

Why it matters: this isn't a chatbot answering chemistry trivia -- it's a model doing a specific, billable step of a trained chemist's job, on real instrument output, fast enough to matter. The gap this closes isn't “can AI discuss science” but “can AI produce a number a lab can sign off on,” and for narrow, well-defined technical tasks like this one, that gap is closing faster than most non-specialists have noticed.

MARKET SIGNAL

AI's capital is consolidating into fewer, bigger bets. SpaceX has now spent \$60 billion on one coding startup and reportedly floated buying a second, rather than build either capability from scratch. OpenAI is sitting on \$122 billion raised in March while its revenue run rate has crossed \$40 billion -- still against roughly \$14 billion in projected annual losses -- and is telegraphing a 2027 IPO. Anthropic, by contrast, just posted its first profitable quarter on a much smaller \$11.5 billion one, and may list even sooner. The two labs are proving that scale and profitability are separate races -- and that going public no longer waits for both to be won at once.

PRACTICAL TAKEAWAYS

1. When you adopt an AI coding tool, track who owns it, not just what it can do.

Cursor changed owners this month, and the startup behind Devin was reportedly a target days later. A tool's roadmap, pricing, and even its willingness to work with rival platforms can shift overnight under new ownership -- factor acquisition risk into any vendor you build a workflow around.

2. Before you rely on an agentic phone or checkout AI, check its disclosure and opt-out rules.

Google's calling agent identifies itself as automated and lets businesses opt out -- that's the bar other agentic-commerce tools should be held to. If you're building or buying one, confirm it discloses itself and gives the other party a way to say no, before you point it at real vendors or customers.

3. Test whether a frontier model can already do a first pass on your specialized technical data.

Claude reading raw NMR and LC-MS files to a lab-grade purity number is a narrow result, but it's a template: pick one well-defined, data-heavy step in your own field and see how close a current model gets before assuming the task still needs a specialist end to end.

4. Read executive comments at all-hands meetings as harder signal than press releases.

Sarah Friar's "2027, or sooner" and her aside about Anthropic's September timeline told employees more about OpenAI's real plans than any official statement has. When you're tracking a company's direction, leaked internal remarks are usually more current than its public messaging.

TOMORROW

Watching whether Cognition says anything further about the SpaceX report beyond Scott Wu's denial, whether Anthropic puts an official date on an IPO to match OpenAI's 2027 target, and whether Anthropic's access program for Claude Science opens up to outside labs.

Topics covered so far: 145 days in, the recurring threads have been agentic AI and autonomy given real-world authority (Claude Code, Cowork, Muse Code, Manus, a real autonomous fighter jet, an AI store manager with hiring and firing power, and now shopping agents that place real phone calls), AI companies converting private momentum into permanent structure (Anthropic's first profitable quarter, OpenAI's \$40 billion run rate and 2027 IPO signal, and now a wave of AI-company acquisitions led by SpaceX's \$60 billion Cursor deal), AI products running into their own unsolved judgment calls (safety-eval gaming, staged capability releases, age-prediction that still can't reliably tell a teenager from an adult), the near-weekly frontier and open-weight release cadence (Gemini, Grok, Qwen, GLM, DeepSeek), AI moving from general chat into specialist professional work (protein design, and now a model reading raw chemistry-lab instrument files to a lab-grade number), security incidents inside AI tooling, and the widening reach of AI regulation and governance (the EU AI Act, UK AISI, Colorado's chatbot law, and labs hiring senior policy officials of their own).