

Anthropic Races Toward the Biggest IPO in History -- Nvidia Quietly Trims a \$105 Billion Bet on OpenAI

Anthropic's bankers could file paperwork within days, a new open payment standard gives AI agents their own wallets, and a validated lab result shows Claude designing working proteins -- three stories about who controls the money in AI's next phase.

Day 147 · Saturday, August 22, 2026

Why this matters: today's stories are all about money changing hands -- who funds the next wave of AI infrastructure, who gets to cash out first, and who pays whom when the AI itself is the one doing the transacting. Anthropic is sprinting toward what could be the largest public offering ever, Nvidia is quietly hedging an infrastructure promise it made five days earlier, and a new payment protocol just gave AI agents the ability to spend real money without asking a human first.

1) Anthropic Could File for the Biggest IPO in History -- As Soon as This Month

Bloomberg reports Anthropic's CFO has told investors the company could file its IPO paperwork before the end of August, targeting a Nasdaq listing around October with Goldman Sachs, JPMorgan, and Morgan Stanley leading an offering that could raise more than \$60 billion. Investors are reportedly underwriting the deal at a valuation near \$2 trillion -- a number that would make it one of the largest public debuts in stock market history, on par with or bigger than SpaceX's record-setting listing earlier this year.

The financial backdrop is moving fast: Anthropic's annualized revenue -- a single month's revenue multiplied by twelve, a snapshot of run-rate, not a guarantee of the year ahead -- crossed \$47 billion in May and an estimated \$65 billion by July. That surge is what let the company report its first-ever operating profit this quarter, two years ahead of its own internal schedule. Going public would make Anthropic the first AI safety-focused lab to open its books to public markets, forcing the same quarterly scrutiny that hits every other public company onto a business whose entire pitch has been long-term caution over short-term growth.

2) Nvidia Scales Back Its Own \$105 Billion Promise to OpenAI -- Days After Making It

On August 17, Nvidia agreed to back up to \$105 billion in financing for a new Ohio data center campus OpenAI plans to lease for twenty years -- about eight gigawatts of capacity, built and owned by SoftBank-backed SB Energy, running exclusively on Nvidia chips. Within days, reports emerged that Nvidia had quietly scaled back parts of that funding guarantee, with the final commitment landing well below early reporting of a \$250 billion figure.

The pullback matters because of what analysts call circular financing: a chipmaker investing in the very customer that buys its chips, which can make both companies' growth numbers look stronger than the underlying demand actually is. Nvidia, OpenAI, and several of the biggest cloud players are now tied together through a web of these deals -- so when the chip supplier itself starts trimming its commitments, it is a signal worth watching. The practical lesson: when you read that a company's revenue or capacity is

"backed" by a partner, check whether that partner is also a major customer -- the growth may be real, or it may just be money moving in a circle.

3) AI Agents Just Got Their Own Wallets -- No Human Approval Required

Cloudflare has expanded its support for x402, an open payment standard built with Coinbase that revives the old, mostly unused HTTP "402 Payment Required" status code. In practice, it lets an AI agent pay a few cents for an API call mid-task -- fetching a dataset, calling a paid model, buying a piece of compute -- without a person clicking "approve" in the loop. More than twenty companies have already adopted it.

The safety rails are the interesting part: a human-owned Account Wallet funds one or more Virtual Wallets that are handed to specific agents, each capped with a maximum total spend, a whitelist of approved merchants, and a limit on any single transaction. That split is what keeps an autonomous agent's budget bounded instead of open-ended. This is the plumbing that turns "agentic AI" from a slideshow buzzword into something that can actually transact in the world -- constrained by rules a human sets in advance, not permission asked in the moment.

VIRAL / OPEN-SOURCE SPOTLIGHT

Claude Designs Working Proteins -- Validated by Labs That Had No Reason to Be Nice About It

Anthropic published lab-validated results this month showing that Claude (Mythos Preview and Opus 4.8) designed protein binders -- molecules built to latch onto a specific target, the basis of most modern drugs -- that worked against 14 of 15 targets tested by two independent contract labs, Adaptic Bio and Twist Bioscience. Success rates landed at 22-35%, roughly double the 10-15% typical for this kind of computational protein design. What's making this spread through AI and biotech circles this week isn't the headline number -- it's who ran the test. Adaptic and Twist are commercial labs paid to synthesize and test candidate proteins for anyone, including Anthropic's competitors, so they had nothing to gain from a favorable result. That is a meaningfully stronger form of evidence than a self-reported benchmark score, and it's why this result is getting cited as a genuine sign that language models are starting to do useful work in wet-lab science, not just describe it.

MARKET SIGNAL

Gartner now predicts more than 40% of agentic AI projects will be cancelled by the end of 2027, citing escalating costs, unclear business value, and weak risk controls -- even as Deloitte finds 74% of organizations still expect to deploy agentic AI within the next two years. Read together, the two numbers describe the same moment from opposite sides: enterprise appetite for AI agents keeps climbing, but a large share of that enthusiasm is being spent on pilots without a defined success metric, a budget ceiling, or a plan for what happens if the agent gets it wrong. The gap between those two numbers is where most of the next two years of agentic AI news is going to come from.

PRACTICAL TAKEAWAYS

1. Read "annualized revenue" numbers as a snapshot, not a forecast

Anthropic's jump from \$47B to \$65B annualized in two months sounds explosive, but it's one month's number stretched across twelve -- useful for spotting trend direction, not for predicting the actual year-end total. Before quoting an "annualized" figure to someone else, say out loud what month it's based on.

2. Before turning on autonomous spending for an AI agent, set the caps first

Whether you're using x402-style agent wallets or any tool-calling agent with a company card attached, insist on a hard spend ceiling, a merchant allowlist, and a per-transaction limit before autonomy is switched on -- exactly the pattern Cloudflare just shipped as a default, not an afterthought.

3. Give every agentic AI pilot a success metric on day one

Gartner's 40% cancellation number tracks projects that drifted without a defined target. If you're greenlighting an agent pilot at work, write down what "working" looks like -- in hours saved, error rate, or dollars -- before the first line of the agent's workflow gets built.

4. Weigh a vendor's AI claim by who verified it, not just what it says

Claude's protein-design result carries more weight because two independent, commercially neutral labs ran the tests -- not Anthropic's own benchmark suite. When a company claims an AI breakthrough, the first question worth asking is: who checked this, and did they have any reason to make it look good?

Topics Covered So Far

146 days in, the recurring threads have been: the compute and chip race (Nvidia, Cerebras, Broadcom, and the data-center financing behind them), a relentless cadence of frontier model releases (GPT-5.x, Gemini 3.x, Claude 4.x, Grok 4.x, and fast-improving open-weight models), the rise of agentic and coding-agent tools (Cursor, Devin, Meta's coding agents), AI safety and governance friction (EU rules, watermarking mandates, safety-exam controversies), corporate finance turning into front-page news (Anthropic's profitability, OpenAI's losses, and now IPO chatter), and AI reaching into the physical and creative world (an Air Force F-16, robotaxis, and AI video tools like Higgsfield). Today's stories -- an IPO, a scaled-back mega-deal, and agents that can spend money on their own -- are the corporate-finance and agentic threads converging in one day.

Tomorrow:

We'll watch whether Anthropic's IPO paperwork actually lands, dig into what the filing would reveal about how the company really makes money, and check whether Nvidia's financing pullback on the OpenAI Ohio deal shows up anywhere else in the AI infrastructure chain.

Generated for varunsingla.com · Daily AI Learning