

Nvidia Made \$96 Billion in Three Months Wall Street Barely Blinked

Plus: Hugging Face is fielding buyout offers near \$13 billion after turning down Nvidia's money earlier this year, and OpenAI built a way to catch AI misuse without ever reading your prompts.

Day 152 · August 27, 2026 · Curated for Varun Singla

Today's throughline is what happens after the hype cycle prices something in. Nvidia delivered the earnings beat everyone expected, and the market's shrug says more about where AI enthusiasm stands than the number itself. Hugging Face, the open-model hub nearly every AI team depends on without thinking about it, is quietly exploring a sale that could hand control of that dependency to a single buyer. And OpenAI shipped a genuinely new answer to a problem that sounded unsolvable: catching misuse without reading anyone's data.

1) Nvidia Made \$96 Billion in Three Months — the Market Priced It In Already

Nvidia's fiscal second-quarter revenue hit \$96.2 billion, more than double the year-ago quarter, with data center sales alone reaching \$89 billion — 93% of everything the company sold. Adjusted earnings of \$2.22 per share beat the \$2.08 analysts expected. Guidance for the current quarter came in at \$108 billion, and CFO Colette Kress told analysts Nvidia now expects \$3 trillion to \$4 trillion in AI infrastructure spending industry-wide by the end of the decade. The stock rose less than 4% in after-hours trading — a modest bump for numbers that would have been stunning three years ago.

Explained simply: an earnings “beat” means a company made more money than Wall Street predicted. A muted stock reaction to a beat usually means investors had already priced in success — they weren't surprised, so there was nothing left to reward. The forward guidance number (\$108 billion) matters more here than the trailing one, because it's Nvidia's own bet on what customers will keep buying, not what they already bought.

Why it matters: this was the earnings report that landed one day after Nvidia laid out its next chip architecture at Hot Chips 2026 — the roadmap slide and the sales numbers, side by side. Both came in strong. The market's flat reaction isn't doubt about AI demand; it's a sign that “Nvidia beats expectations” has become the base case rather than the surprise, which is its own kind of signal about how far this cycle

has run.

2) Hugging Face Is Exploring a Sale That Could Value It at \$13 Billion

Hugging Face — the platform where developers and researchers upload, share and download AI models and datasets, used by nearly every team building on open-weight models — is reportedly fielding acquisition offers that would value it near \$13 billion. No deal has been reached, and it isn't public who's bidding, but the company has been talking to banks to evaluate offers. Hugging Face last raised money in 2023 at a \$4.5 billion valuation, and earlier this year turned down a \$500 million investment from Nvidia that would have valued it at \$7 billion. Co-founder Clément Delangue has spoken publicly about a “long-term responsibility” to the community that trusts the platform with its models and data.

Explained simply: think of Hugging Face as the GitHub of AI models — a shared, mostly neutral place where anyone can publish a model and anyone else can download it, with no single company controlling who gets access. That neutrality is the entire value proposition. If one company — especially a frontier lab or cloud provider that also competes with the models hosted there — ends up owning it, neutrality is no longer guaranteed by design; it becomes a policy choice the new owner makes.

Why it matters: open-weight AI's whole pitch is that no single company gets to decide who can build with a given model. That pitch depends on the distribution layer staying neutral. A sale doesn't have to go badly to still change the calculus — even a well-intentioned buyer changes who gets to make that call, and offers nearly tripling the valuation Hugging Face rejected earlier this year show how much strategic value someone now sees in owning that leverage point.

3) OpenAI Built a Way to Catch Misuse Without Reading Your Prompts

OpenAI announced Private Safety Processing, a system that watches for patterns of AI misuse across multiple conversations while keeping its Zero Data Retention policy fully intact — meaning prompts and responses are still discarded after each request and never seen by OpenAI staff. When the automated system flags a pattern, OpenAI receives only a narrow category label describing the type of concern, not the content that triggered it. The feature is in early testing with a small group of enterprise customers, including Glean, Databricks, Abridge and Microsoft, with a wider rollout and technical white paper due in September. The move directly answers Anthropic's newer policy of retaining enterprise data for 30 days on covered models for safety purposes — the two companies are now explicitly competing on privacy architecture, not just model quality.

Explained simply: safety monitoring and zero data retention have always pulled in opposite directions — you can't spot a pattern across ten conversations if you're deleting each one right after it happens. OpenAI's answer is to keep an encrypted, automated signal about what kind of behavior looks concerning, without keeping the behavior itself. It's the difference between a smoke detector that reports “smoke detected in unit 4B” and a security camera that records everything happening inside unit 4B.

Why it matters: enterprise buyers have been forced to choose between “trust us to keep you safe” and “trust us to never look.” If Private Safety Processing works as described, that's no longer a strict tradeoff — which is exactly the kind of technical answer that turns a policy debate into a procurement checkbox, and puts real pressure on every other AI vendor's retention policy to explain why it needs to see more.

VIRAL / OPEN-SOURCE SPOTLIGHT

We Had ChatGPT Make Us a Script

The viral trend turning group chats into deadpan theater. The mechanic is dead simple: a group of friends prompts ChatGPT for a short scene — a breakup, an argument, a wildly specific hypothetical — then films themselves performing the AI's dialogue completely straight-faced, no matter how unhinged the lines get. A caption up front (“we had ChatGPT make us a script”) tells viewers exactly what they're watching before the first line lands.

Why it's taking off: the comedy comes from two places at once — how strange AI-generated dialogue sounds once a real person has to say it out loud, and how committed the group stays even when the scene makes no sense. It costs nothing to make, needs no editing skill, and rewards exactly the kind of chaotic group-chat energy that already exists among friends — which is why it's spreading through group chats faster than most produced content can.

MARKET SIGNAL

Hugging Face's implied valuation has nearly tripled in eight months — from the \$7 billion it turned down from Nvidia earlier this year to the roughly \$13 billion now on the table. That jump isn't about the platform doing three times the business; it's about how much more a buyer will now pay for control of the open-model distribution layer itself, on the same week Nvidia told investors to expect \$3 trillion to \$4 trillion in AI infrastructure spending by 2030. The money is flowing toward chokepoints — chips and distribution — as much as it's flowing toward the models themselves.

PRACTICAL TAKEAWAYS

1. If your stack depends on Hugging Face for model hosting, plan for a possible ownership change now.

A sale near \$13 billion looks increasingly likely — start asking what a change of owner would mean for the models and datasets your team relies on there, before it happens rather than after.

2. Read Nvidia's guidance number, not the beat, if you're forecasting AI infrastructure costs.

The \$108 billion Q3 guide and the \$3–\$4 trillion decade-end forecast are Nvidia's own bet on sustained demand — more useful for planning than a trailing quarter you already knew would be strong.

3. If you're evaluating an AI vendor on privacy, ask specifically how it detects misuse without storing your data.

OpenAI's category-only signal design is now the concrete answer to a tradeoff every other vendor has been hand-waving around — use it as the bar for what “we take privacy seriously” should actually mean.

4. Try the “AI script” format as a genuinely low-effort content or icebreaker exercise.

Constrain the prompt tightly, then deliver the output completely straight — the gap between AI weirdness and human commitment is doing all the work, and it costs nothing to test.

TOMORROW

With August closing out, the next question is whether Anthropic actually files paperwork for what people close to the process say could rival SpaceX's IPO in size. Tomorrow: watching for the filing itself, not just more reporting about one.

Topics Covered So Far

Foundations: multi-agent orchestration, MCP and A2A protocols, plan-and-execute patterns, long-horizon task capability.

Frontier models & infrastructure: coding-agent wars (Meta, Cognition), open-weight model churn, Nvidia chip economics and the Vera/Groq roadmap now backed by a \$96 billion quarter, GLM-5.3's emergent cyber capability, OpenAI Astra's formally verified math proofs.

Business & policy: Anthropic's first operating profit and its approaching IPO, OpenAI's revenue and ChatGPT-for-teens rollout, Nvidia trimming its OpenAI bet, Google DeepMind's leadership reset, the EU AI Act's transparency rules taking effect, and now a possible \$13 billion sale of Hugging Face.

Agentic risk & infrastructure: AI store managers replacing human staff, autonomous agent wallets and their first attack research, summer 2026 safety-eval containment failures, an AI assistant disclosing its own vulnerability under questioning, a sandbox escape its maker chose not to patch, and OpenAI's new answer to monitoring misuse without storing data.