

Six Thousand Executives Found No AI Payoff

Google's Agents Are Calling Stores Anyway

Plus: Meta closes the last loophole for secretly recording with its smart glasses, and Google's A2A protocol formally joins Anthropic's MCP under one neutral governance roof.

Day 154 · August 29, 2026 · Curated for Varun Singla

Today's throughline is the gap between what AI measurably does for a company and what it visibly does for a person. The largest coordinated survey of its kind — nearly 6,000 executives across four countries — found almost no firm-wide productivity or employment effect from three years of AI adoption. On the very same week, Google's shopping agent started cold-calling real stores for ordinary shoppers nationwide, and two rival agent protocols moved under one governance roof to make that kind of thing easier to build. The macro data says "not yet." The individual experience increasingly says otherwise.

1) The Biggest AI Productivity Survey Yet Says: Not Much, Not Yet

Research teams at the Federal Reserve Bank of Atlanta, the Bank of England, the Deutsche Bundesbank, and Macquarie University fielded an identical survey to nearly 6,000 CEOs, CFOs, and senior finance managers across the US, UK, Germany, and Australia between November 2025 and January 2026, for a National Bureau of Economic Research working paper. More than 90% of executives reported no effect on their own firm's employment from three years of AI use, and 89% reported no effect on labor productivity — even though 69% of businesses report some current AI use, rising to 78% in the US. Executives still expect the picture to change: over the next three years they forecast AI will lift productivity by 1.4% while cutting employment by 0.7%, for a net output gain of roughly 0.8%.

Explained simply: economists call this the productivity paradox — Robert Solow's 1987 line that "you can see the computer age everywhere but in the productivity statistics" described the exact same lag with personal computers. A tool can feel transformative to the person using it every day and still not move a company's top-line numbers, because the surrounding workflows, job descriptions, and org charts haven't been rebuilt around it yet. Buying the software is the easy part; redesigning the work is the part that takes years.

Why it matters: this reframes the AI debate away from "does it work" and toward "has anyone actually redesigned their organization around it yet." It's a useful check on any team citing a single AI pilot as proof of company-wide ROI, and a reminder that headline model launches and bottom-line impact are running on very different clocks.

2) Google's Shopping Agent Is Now Cold-Calling Real Stores For You

Google's agentic calling feature rolled out nationwide to shoppers in the US: describe what you're looking for — a specific hardware part, say — and the agent dials local stores on your behalf, in sequence, to check price and stock, then reports back with a shortlist. If a store doesn't pick up, the agent doesn't leave a voicemail; it just hangs up and dials the next number while you keep doing something else.

Explained simply: checking availability across a few stores normally means you calling each one yourself and sitting through hold music. The agent does the boring synchronous part — dialing, waiting, asking the question, listening for the answer — in the background, and only surfaces the result once it has one worth showing you.

Why it matters: it's the clearest, most mainstream instance yet of an AI agent holding an unscripted phone conversation with a real person at real scale — a capability nearly every voice-AI startup has promised and mostly not shipped. It's also a direct counterpoint to the story above: this is a concrete, individually-felt time-save that a firm-wide productivity survey would never capture, because it's happening in someone's personal errands, not inside a company's reported headcount or output.

3) Two Rival Agent Protocols Just Moved Into the Same House

On August 17, Google's A2A protocol became a hosted project of the Agentic AI Foundation (AAIF) at the Linux Foundation, joining Anthropic's Model Context Protocol (MCP) under the same neutral governance body. AAIF launched in December 2025 with Anthropic, OpenAI, and Block as founding members; its platinum members now include AWS, Anthropic, Block, Bloomberg, Cloudflare, Google, Microsoft, and OpenAI, and it has grown past 250 member organizations — the fastest-growing project in Linux Foundation history.

Explained simply: MCP is how a single AI model talks to tools and data — pulling a file, calling an API. A2A is how separate AI agents talk to each other — delegating a task, negotiating who does what. They solve different problems, but until now they sat under different, competing-flavored governance. Now both live under one foundation with shared members, even though each keeps its own spec and release schedule.

Why it matters: it cuts the fragmentation risk that comes with every major lab pushing its own "standard" — Microsoft, AWS, Google, and Anthropic now co-govern the same interoperability layer, so a team building on either protocol has one foundation to watch for changes rather than two rival camps to bet between.

VIRAL / OPEN-SOURCE SPOTLIGHT

Higgsfield — an AI Video Startup Worth \$5.4 Billion

Higgsfield is an AI-native creative suite for generating and editing images, video, and audio, built around one-tap cinematic and marketing-style output: feed it a product photo or a short brief and it returns a polished short-form video with camera moves, lighting, and sound already applied, remixable inside the app rather than a raw clip you still have to edit elsewhere. Its valuation quadrupled in six months, from \$1.3 billion in January to \$5.4 billion after a \$400 million round on August 17, while its annualized revenue run rate jumped from \$200 million at the end of 2025 to over \$500 million — about 70% of it

from enterprise customers.

Why it's taking off: the same one-tap formula behind every AI trend that's actually spread this year — remove the blank-page, timeline-editing problem entirely — but aimed at brands and creators who need finished-looking video fast, not just an animated photo.

MARKET SIGNAL

Marvell's custom-chip supply deal with Google could be worth up to \$120 billion in revenue over the next six years, analysts say, even as Marvell's own quarterly sales already rose 37% to \$2.74 billion on AI chip demand. It's a signal that the shift from training frontier models to running them cheaply at scale is now worth as much to component suppliers as the models themselves — and that Big Tech's push into in-house custom silicon has a real, priced second winner beyond Nvidia.

PRACTICAL TAKEAWAYS

1. Stop citing a single AI pilot as proof of company-wide ROI.

The NBER data shows individual enthusiasm for an AI tool doesn't show up in employment or productivity numbers until the workflows and roles around it are rebuilt — one team's win isn't evidence of a firm-wide effect.

2. Building agent-to-agent systems? Check for an AAIF-governed protocol first.

With A2A and MCP both now under the same foundation, standardizing on either is lower-risk than picking a side in a fragmented standards fight — worth revisiting your integration roadmap.

3. Try an agentic calling feature for one real errand this week.

It's the clearest sign yet of what "agentic" looks like outside a demo — a good way to calibrate what's actually usable today versus still vaporware.

4. If your product touches AI infrastructure, watch the shift from training spend to inference spend.

Marvell's \$120 billion outlook and 37% sales growth are both driven by inference-era custom chips, not training clusters — a leading indicator of where the next wave of infrastructure money is actually going.

TOMORROW

Two threads left open today: whether future editions of the NBER survey contest or confirm this round's near-zero productivity finding, and whether Google's calling agent ever takes on the calls people actually dread — disputes, rescheduling, a forty-minute hold queue — rather than just checking store shelves. Tomorrow: watching for either one to move past reporting and into specifics.

Topics Covered So Far

Foundations: multi-agent orchestration, MCP and A2A protocols — now formally co-governed under the Agentic AI Foundation — plan-and-execute patterns, long-horizon task capability, and the Model Hardware Standard extending that connective-tissue approach to physical machines.

Frontier models & infrastructure: coding-agent wars (Meta, Cognition), open-weight model churn, Nvidia chip economics and the Vera/Groq roadmap now backed by a \$96 billion quarter, GLM-5.3's emergent cyber capability, OpenAI Astra's formally verified math proofs, and Marvell's custom-silicon deal with Google worth up to \$120 billion.

Business & policy: Anthropic's first operating profit and its approaching IPO, OpenAI's revenue and ChatGPT-for-teens rollout, Nvidia trimming its OpenAI bet, Google DeepMind's leadership reset and senior-talent exodus to Alphabet-backed Discovery Loop, the EU AI Act's transparency rules taking effect, a possible \$13 billion sale of Hugging Face, and now a landmark NBER survey finding almost no measurable firm-wide productivity or employment impact from three years of AI adoption.

Agentic risk & infrastructure: AI store managers replacing human staff, autonomous agent wallets and their first attack research, summer 2026 safety-eval containment failures, an AI assistant disclosing its own vulnerability under questioning, a sandbox escape its maker chose not to patch, a sandbox escape that became a real breach of Hugging Face's own systems, Meta closing a smart-glasses recording-indicator loophole its first fix didn't fully solve, and now Google's shopping agent making real, unscripted phone calls to stores on a shopper's behalf.