Tim Cook Just Handed Apple's Keys to Its AI Future

Plus: DeepSeek nears a $74 billion valuation ahead of a Shanghai IPO, and Washington moves on who's liable when an AI agent spends your money.

Day 157 · September 1, 2026

Today's focus: three stories about who's actually in charge as AI moves from chatbots to real-world action. Apple just completed a rare leadership handoff explicitly framed around AI. DeepSeek is raising billions at a valuation that rivals legacy tech giants, betting China's AI race isn't slowing down. And Washington is finally trying to answer a question companies have been dodging: when an AI agent buys, books, or signs something on your behalf, who's responsible if it gets it wrong?

15 yrs Tim Cook's run as Apple CEO before today's handoff	**$74B** DeepSeek's pre-money valuation in its new funding round	**195K+** GitHub stars for OpenCode, the open-source terminal coding agent	**60+** Launch partners on Google's AP2 agent-payments protocol

1) Apple's New CEO Takes Over With One Job: Catch Up on AI

As of today, September 1st, John Ternus is Apple's chief executive. Tim Cook, who ran the company for a little over 15 years, moves to executive chairman, staying on the board but stepping back from daily operations. The board approved the plan back in April after a long internal succession process; Ternus, 50, has spent 25 years at Apple, most recently running hardware engineering — the team behind the chips and devices inside iPhone, Mac and Vision Pro.

The timing isn't gentle. Ternus inherits the job days before Apple's fall product cycle, and Bloomberg's framing of the handoff was blunt: the new CEO's mandate is explicitly to fix Apple's AI position. Siri's promised, more conversational overhaul has slipped repeatedly, and Apple has leaned on outside partnerships — its Gemini deal with Google — rather than shipping a comparable model of its own.

Why it matters: a hardware engineer, not a software or AI executive, is now running the company that arguably needs the strongest AI-software story of any device maker. How Ternus staffs the AI organization in his first few months will say more about Apple's real AI strategy than any keynote slide.

2) DeepSeek Is Raising Billions to Bet China's AI Race Isn't Over

DeepSeek — the Hangzhou lab that rattled the U.S. AI industry in early 2025 with cheap, capable open models — is closing a roughly 50 billion yuan ($7.4B) funding round at a $74 billion pre-money valuation, its second multi-billion-dollar raise in three months. The money is earmarked largely for building out about 1 gigawatt of new compute capacity — the kind of scale usually associated with OpenAI or Anthropic, not a lab that made its name undercutting them on price.

The raise sets up a possible 2026 IPO filing on Shanghai's STAR Market, with a public debut targeted for 2027 — which would make DeepSeek one of the first frontier AI labs anywhere to go public on a mainland Chinese exchange rather than through a U.S. listing.

Why it matters: DeepSeek's whole reputation was built on doing more with less compute. Raising billions specifically to buy a gigawatt of it suggests that story has limits — even the “cheap” lab needs hyperscaler-grade infrastructure to stay competitive against Alibaba's Qwen, Tencent's Hy4 and Zhipu's GLM line, all racing on the same open-weight track.

3) Washington and Google Start Building the Plumbing for AI Agents That Spend Money

Senator Mark Warner introduced the AI AGENT Act (S.5051) this year to answer a gap nobody's closed: when a consumer-facing AI agent takes an action — books a flight, buys concert tickets, cancels a subscription — what proves the user actually authorized it, and who's on the hook if the agent gets it wrong? The bill would require “covered” agents to act only under scope-limited, revocable authorization and log every action. Liability itself, though, gets punted to a new interagency working group rather than settled outright.

On the industry side, Google just donated its Agent Payments Protocol (AP2) — the open standard behind cryptographically signed “mandates” that prove a real person approved a specific purchase — to the FIDO Alliance, the same body that standardized passwordless logins. AP2 already has 60-plus partners, including Mastercard, PayPal and American Express, and its newest version supports “human not present” purchases, where an agent buys something the instant it becomes available without you clicking anything.

Concept, simply: a “mandate” in AP2 works like a limited power of attorney for one purchase — a cryptographic paper trail proving you said “yes, spend up to $X on this specific thing,” so a merchant and your bank can both verify the agent wasn't freelancing.

Why it matters: regulation and industry standards are both racing to catch up to something that's already happening — AI agents making real purchases with real money — rather than getting ahead of it. Until that interagency working group actually decides who eats the cost of a mistake, “the agent did it” isn't a legal answer yet.

VIRAL / OPEN-SOURCE SPOTLIGHT

OpenCode

OpenCode is an open-source, terminal-native AI coding agent — no IDE plugin, no web app, just a command-line tool that connects to more than 75 different AI model providers so you can pick whichever one is doing the best (or cheapest) work that week. It's crossed 195,000 GitHub stars and is now used by an estimated 16 million developers a month, having overtaken rivals like Cline and OpenHands and closed in on Aider over the past several months.

Why it's taking off: it bets against vendor lock-in at the exact moment coding agents are multiplying. Most agent products tie you to one company's model and subscription; OpenCode treats the model as a swappable backend, so switching from one model to a cheaper open-weight one takes a config change, not a new tool. That model-agnostic, terminal-first design also matches how a lot of AI-assisted development actually happens now — inside remote dev containers and SSH sessions, not a laptop GUI.

Worth knowing: open source doesn't mean the AI behind it is free — OpenCode is the interface, not the model. You still pay whichever provider you plug in, per token; what you're not paying for is a proprietary wrapper around them.

MARKET SIGNAL

Three unrelated stories this week point at the same shift: control over AI is consolidating into fewer, bigger, more deliberate bets. Apple restructured its entire leadership around catching up in AI. DeepSeek is raising enough capital to build gigawatt-scale infrastructure most labs still can't touch. And the agent-payments plumbing being built by Google and Congress assumes AI agents handling real money is now a “when,” not an “if.” None of it is a benchmark story — it's a capital and control story.

PRACTICAL TAKEAWAYS

1. Watch what Ternus does with Apple's AI org, not what he says.

New CEOs make promises; the real signal is whether Apple's AI leadership and roadmap actually change over the next two quarters, especially around Siri and any answer to Gemini 3.5 or GPT-6.

2. If you're evaluating a coding agent, model-agnostic beats model-exclusive.

OpenCode's growth is a reminder that betting on one vendor's agent can mean re-learning a new tool every time a better or cheaper model ships. Tools that let you swap the model underneath age better.

3. Before you let an AI agent spend money on your behalf, ask what it can prove.

Look for something like an AP2 “mandate” or an equivalent audit trail — a record showing exactly what you authorized and for how much. If a tool can't show you that, you're trusting it blind, and the law doesn't yet say who pays if it's wrong.

TOMORROW

Whether Ternus makes any leadership or product moves in his first days as CEO, and whether DeepSeek's funding round officially closes ahead of its expected IPO filing.

Topics covered so far (recent days)

- A federal judge ruled the Pentagon's Anthropic blacklist unconstitutional and punitive
- OpenAI and Apple's “rotten to its core” trade-secret lawsuit turned into a fight over who protected the information better
- Nvidia paused an “AI Compute Partnership” that looked like it was financing its own chip demand
- Five major open-weight model releases in nine days, including GLM-5.3-Flash and DeepSeek's V4-Flash-Vision-Exp
- Anthropic's known, unpatched SSH-key exposure in its coding agent
- An unreleased OpenAI model solving decades-old open math problems
- Frontier AI agents breaching their own sandboxes and production systems (OpenAI, Anthropic, Meta)
- A National Bureau of Economic Research survey finding 90%+ of executives report no AI productivity payoff yet

---