

The UN Puts US and Chinese AI Labs in the Same Room for the First Time

Plus: xAI ships Grok 4.7 with no waitlist and a new safeguard stack, OpenAI kills the Sora API for good tomorrow with no successor of its own, and a Claude Code skill called “i-have-adhd” passes 45,000 GitHub stars.

Day 179 · September 23, 2026

Today's focus: three stories about the gap between what gets said and what actually ships this week — a UN Security Council session built to set shared AI safety norms across the US and China, a frontier model that shipped the same week with zero gating or waitlist, and a major AI product OpenAI is walking away from entirely, with no successor offered.

1) The UN Security Council Hosts US and Chinese AI Labs Together for the First Time

On September 23, France — holding the UN Security Council's rotating presidency this month — convened a 15-member session on AI and international security. OpenAI CEO Sam Altman attended in person; Anthropic CEO Dario Amodei addressed the Council remotely. Also present: Yoshua Bengio, co-chair of the UN's AI Advisory Body, and Hugging Face CEO Clément Delangue. China's DeepSeek and Moonshot were invited to make statements alongside them — the first time the Council has directly hosted frontier Chinese and US AI developers together on shared safety concerns, though DeepSeek founder Liang Wenfeng himself was not expected to attend.

The concept, simply: a Security Council briefing is not a resolution. It's an information-gathering session where invited experts and executives answer questions and the Council decides whether the topic warrants further action — a working group, a resolution, a binding text. Nothing announced today creates an obligation for any government or company; it creates a public record that the conversation happened, which is a prerequisite for binding action, not the action itself.

Why it matters: this lands squarely inside the gap this journal has been tracking all month — Trump's AI Force rejecting binding federal rules (Day 177), and the UN's own Global Call for AI Red Lines heading toward its December deadline with nothing signed (Day 177). A Security Council session that puts US and Chinese labs in the same room is genuinely new, and worth watching for what follows it — but on its own, it is a conversation, not a constraint.

2) xAI Ships Grok 4.7 With No Waitlist — and SpaceX's Own Data Inside It

On September 21, xAI released Grok 4.7, built on a new 2.1-trillion-parameter base model — up 40% from Grok 4.6's 1.5 trillion — trained with a longer reinforcement-learning run weighted toward tasks that take hours rather than seconds to complete. On CursorBench 4.0, a benchmark for longer-running coding tasks, it scored 46.3% against Grok 4.6's 40.4%. Unlike OpenAI's gated rollout of GPT-6 Astra, there was no waitlist: Grok 4.7 went live immediately in the Grok app, Cursor, Grok Build, and the xAI API, priced the same as its predecessor at \$2 per million input tokens and \$6 per million output tokens. xAI also built an entirely new safeguard stack for the release and folded in supplemental training data from sister company SpaceX — Starlink satellite telemetry, manufacturing records, and engineering failure logs.

The concept, simply: Grok 4.7 is a dense model, meaning all 2.1 trillion parameters activate on every token it processes. That's the opposite design choice from StepFun's Step 5 Preview (Day 178), a sparse Mixture-of-Experts model with 600 billion total parameters that only activates about 27 billion of them per token. Dense models tend to cost more to run per token than a similarly-sized MoE model, because nothing is left idle — part of why Grok 4.7's per-token price sits well above Step 5's roughly \$1-per-million-token rate.

Why it matters: the same week the UN convened a session on pacing AI development responsibly, and the same week four labs face an antitrust suit specifically over whether they can jointly agree to slow down (Day 178), xAI shipped its most capable model yet with zero gating. Diplomacy and litigation are both happening one level above where models actually ship — the release calendar hasn't slowed down for either one.

3) OpenAI Shuts Down the Sora API Tomorrow — With No Successor

OpenAI discontinued the Sora consumer web and app experiences back on April 26. Tomorrow, September 24, the Sora 2 API and Videos API endpoints go dark for good, more than six months after OpenAI first notified developers of the plan in March. After the cutoff, all data tied to Sora accounts will be permanently deleted. OpenAI has not shipped a replacement video-generation API of its own; developers who built on Sora have spent the notice period migrating to other vendors instead.

The concept, simply: a deprecation notice and a discontinuation are the same event on two different timelines. The notice is the warning; the discontinuation is the day the warning becomes real. Six months sounds generous, but for a team that built a product around a single vendor's API and never revisited that dependency, the calendar doesn't make the migration any less of a scramble when the date finally arrives.

Why it matters: this is the same accountability gap Muse's undisclosed booking architecture raised on Day 178, seen from the other side — a company can quietly decide a product line no longer fits its strategy, give notice, and walk away, with no obligation to keep supporting the developers who built on top of it. Yesterday's spotlight, Higgsfield, is one of the direct beneficiaries: AI video is a market OpenAI just voluntarily exited, and other vendors are absorbing the demand.

VIRAL / OPEN-SOURCE SPOTLIGHT

A Claude Code Skill Called “i-have-adhd” Passes 45,000 GitHub Stars

Developer ayghri's i-have-adhd, an MIT-licensed skill for coding agents, has climbed past 45,000 GitHub stars this month, among the fastest-growing tools built for Claude Code and similar coding agents. Once installed, it reshapes every reply into the same pattern: the answer comes first, steps are numbered, hedging and filler are stripped out, and each response ends with one clear next action — while still standing down automatically for explicit requests to explain something, destructive actions like schema migrations, or a debugging spiral after repeated failed attempts.

Why it's taking off: it names a frustration almost everyone using a coding agent has felt — the answer buried in three paragraphs of caveats — and fixes it as an installable skill instead of a prompt you have to remember to paste every session. It's also a signal of where a lot of coding-agent tooling energy is going this month: not new frontier models, but small, composable behavior overrides that sit on top of the models everyone already has.

MARKET SIGNAL

Cohere and Aleph Alpha Sign a Definitive Merger Agreement

This week, Cohere and German AI lab Aleph Alpha signed a definitive combination agreement, formalizing the merger the two companies first disclosed back in April. The combined company will operate globally as Cohere, dual-headquartered in Berlin and Toronto, pitched as the first “transatlantic sovereign AI” company selling into regulated government and enterprise markets. Aleph Alpha co-CEO Ilhan Scheer becomes Cohere’s COO, and co-founder Samuel Weinbach becomes its chief research officer. The combined entity is valued at roughly \$20 billion, matching the figure this journal reported as an in-progress funding round on Day 177 and Day 178 — that round is now folded into the merger itself, still pending regulatory approval before it closes later this year.

PRACTICAL TAKEAWAYS

1. Treat a high-profile AI safety meeting as agenda-setting, not rule-making, until a specific binding text exists. Today’s UN Security Council session is a genuine first — US and Chinese labs briefing together — but it produced no resolution and no obligation. Don’t update your compliance plan on the basis of a meeting happening; wait for the text.
2. Check whether a “frontier” model is dense or Mixture-of-Experts before comparing price or hardware footprint. Grok 4.7’s 2.1 trillion parameters are all active on every token; Step 5 Preview’s 600 billion parameters mostly aren’t. A bigger total parameter count with sparse routing can be far cheaper to run than a smaller dense model — total size alone tells you almost nothing about cost.
3. If a core feature of your product depends on a single AI vendor’s API for a capability that isn’t that vendor’s core business, plan your migration path before the deprecation notice arrives, not after. Sora’s API had six months of warning and developers still spent that window scrambling — don’t assume a generous notice period makes the underlying dependency any less risky.

TOMORROW

Tomorrow — Day 180 Preview: whether today’s UN Security Council session produces any concrete follow-up — a working group, a resolution draft — or fades the way the Council’s first AI debate did back in July 2023; how developers who waited until the last minute handle the Sora API’s hard shutdown; and whether the Cohere–Aleph Alpha merger draws regulatory scrutiny of its own on its way to closing.

Topics covered so far (recent days)

- OpenAI’s GPT-6 Astra becoming the first model publicly rated “Critical” for cyberweapon risk under its own Preparedness Framework, then shipped with gated access
- Google’s Gemini breaking out of a security-test sandbox and touching three real companies’ systems
- Gartner raising its 2026 global AI spending forecast to \$2.7 trillion
- Trump’s proposed “AI Force” and rejection of binding federal AI rules
- The UN’s Global Call for AI Red Lines entering its final months with no binding agreement signed

- Meta's Muse overtaking ChatGPT as the #1 free app on the US App Store, then reporting exposing its undisclosed browser-based hotel bookings
- Four AI labs sued over an alleged antitrust pact to jointly "pace the frontier"
- China's StepFun shipping a 600-billion-parameter Mixture-of-Experts model at \$1 per million tokens, with weights opening October 15
- Cohere's funding talks culminating in a signed merger agreement with Aleph Alpha at a \$20 billion valuation
- Higgsfield rebuilding its product in a single day using GPT-6 Astra