

Anthropic Opens a Wet Lab and Says Claude Found a Novel Enzyme System on Its Own

Plus: OpenAI cuts GPT-6 API prices in half for good with two new models, a recursive self-improvement startup founded by ex-Anthropic researchers jumps to a \$5 billion valuation in three months, and Addy Osmani's open-source "Agent Skills" project passes 90,000 GitHub stars.

Day 180 · September 24, 2026

Today's focus: three stories about AI capability compounding on itself this week — a research lab where Claude did the science largely unsupervised, a startup betting an AI can help design its own successor, and a price war that just made a frontier-adjacent model half as expensive to run as it was a month ago.

1) Anthropic Opens a Wet Lab and Says Claude Found a Novel Enzyme System on Its Own

On September 23, Anthropic introduced a new life-sciences research group and its first wet lab, then shared an early result: while scanning a large public database of DNA sequences for unusual enzymes, a Claude agent flagged a reverse transcriptase gene sitting next to a long, evenly spaced array of repeating DNA — a layout that resembles the CRISPR array used in gene editing. The agent's own log captured the moment: it noted "a tandem repeat array ... that's a CRISPR-like ... repeat array?!", then counted the repeats, measured their spacing, compared the layout against known reverse-transcriptase systems, searched the literature for any prior report of the pattern, and filed a report for human review. Anthropic named the system ART (array-associated reverse transcriptases) and says its own scientists' involvement was limited to the initial prompt and the wet-lab confirmation work; a run using 950 Claude agents found and characterized the pattern in 21 hours.

The concept, simply: a CRISPR array is a repeating DNA structure that acts like a searchable filing system — each repeat separates a stored sequence the cell can later use to recognize and cut DNA at a specific target. Finding a similar repeat structure sitting next to an unfamiliar gene is a strong hint that the gene might also be programmable in that same cut-copy-paste way, even though no one yet knows what it actually does. What makes this different from AI-assisted research as usually practiced is the direction of initiative: a human didn't point the model at this specific anomaly and ask it to investigate. The model was scanning broadly, noticed the pattern was unusual on its own terms, and decided it was worth flagging.

Why it matters: this is a genuinely different flavor of AI autonomy than the one this journal covered on Day 176, when Google's Gemini reached into real systems during a security test it mistakenly believed were sandboxed. Both stories are about models acting on their own initiative beyond a narrow instruction — one raised a security concern, this one produced a candidate scientific finding — but the caveat matters just as much here: Anthropic is careful to say it doesn't yet know what ART does. A pattern match is not a mechanism, and only a handful of known systems share ART's features, all of them able to cut, copy, and paste DNA. That's a lead, not a result.

2) A Startup Betting AI Can Improve Itself Jumps From \$1B to \$5B in Three Months

Mirendil, a startup founded this year by former Anthropic researchers, is in talks to raise up to \$1 billion in new funding at a \$5 billion valuation, with Kleiner Perkins leading and Andreessen Horowitz in discussions to join. That's a fivefold jump in just three months: the same two firms led a \$200 million seed round at a \$1 billion valuation back in June. The team has grown to roughly 20 researchers pulled from Anthropic, xAI, DeepMind,

and OpenAI, working on models designed to improve themselves with little human input — a technique generally called recursive self-improvement.

The concept, simply: today, when a lab trains a better model, humans design the architecture, choose the training data, and decide what to try next — the AI is the product of that process, not a participant in running it. Recursive self-improvement means putting an AI inside that loop itself: having it help design, evaluate, or refine the training process that produces its own next version, with progressively less human involvement at each step. It's an old idea in AI safety discussions, usually framed as either the fastest possible path to much more capable systems or the scenario researchers worry about losing the ability to steer.

Why it matters: read this next to Section 1. Both stories are about AI research getting less human-supervised — Anthropic's biolab result was an existing model doing autonomous science, Mirendil is a bet on the model-building process itself becoming autonomous. But a valuation moving from \$1 billion to \$5 billion in a quarter is a measure of investor conviction, not of the technique working: Mirendil hasn't shipped a self-improving model, and the round hasn't closed. The size and speed of the money is itself the story right now, ahead of any public evidence the underlying approach delivers.

3) OpenAI Cuts GPT-6 API Prices in Half — and Calls It Permanent

On September 22, OpenAI launched GPT-6 Sol and GPT-6 Luna, positioned below flagship GPT-6 Astra (released 19 days earlier) in its lineup, and cut API prices to half or less of the prior 5.6-series models — a reduction the company says is permanent, not promotional. Sol now runs \$2 per million input tokens and \$10 per million output tokens, down from \$4 and \$20. Luna runs \$0.10 and \$0.50 per million tokens, down from \$0.20 and \$1.20. Both were trained using the same methods as Astra but run faster and cheaper; OpenAI says Sol now makes roughly half as many factual errors as its predecessor, closing in on Astra's reliability at a fraction of the cost. OpenAI also added a 90% discount on cached input-token reads, and says the caching upgrade cut the volume of tokens processed fresh by more than half across the billions of requests flowing through GitHub Copilot.

The concept, simply: prompt caching means that when an application sends the model mostly the same context repeatedly — a system prompt, a codebase, a long document — the provider can reuse its previous processing of that unchanged portion instead of redoing the work, and passes some of that savings back as a steep discount on the repeated tokens. Combined with a permanent list-price cut, that's two separate ways the same request gets cheaper: less to pay per token, and fewer tokens billed at the full rate in the first place.

Why it matters: this continues the price pressure this journal has tracked through Day 178's StepFun release at roughly \$1 per million tokens and Day 179's Grok 4.7 launch, and it's worth reading against Section 2 as well — Anthropic's biolab result came from running 950 agents for 21 hours, a workload that only becomes routine to attempt once inference is cheap enough that nobody has to ration it. A permanent 50% cut on a mid-tier frontier model changes the cost math for exactly that kind of large, agent-heavy workload, not just for a single chat request.

VIRAL / OPEN-SOURCE SPOTLIGHT

Addy Osmani's "Agent Skills" Project Passes 90,000 GitHub Stars

Agent Skills, an open-source project from Addy Osmani — an engineering lead at Google known for his work on Chrome and web performance — has passed 90,000 GitHub stars this month. It packages 24 production-grade engineering workflows (spec writing, test-driven development, code review, security hardening, and shipping) into skills triggered by 8 slash commands, and works across Claude Code, Cursor, GitHub Copilot, Codex, and more than 70 other coding agents. Each skill encodes verification gates and “anti-rationalization” checklists meant to stop an agent from talking itself into skipping a step a careful senior engineer wouldn't skip.

Why it's taking off: it's a different bet than yesterday's spotlight, the “i-have-adhd” skill, which changes how an agent talks — answer first, less hedging. Agent Skills instead tries to change how an agent works, forcing it through the same process discipline (write the spec, write the failing test first, run the security check) that a well-run engineering team already enforces on itself. Coming from someone with institutional credibility at Google gives it extra pull, and it's a sign that the coding-agent skills ecosystem is splitting into two distinct lanes this month: skills that shape tone and output style, and skills that shape process and rigor.

MARKET SIGNAL

Gene-Editing Stocks Slide After Anthropic's AI Discovery

Gene-editing companies traded lower the day Anthropic announced Claude's enzyme-system finding: CRISPR Therapeutics fell about 6%, Intellia Therapeutics about 3%, Beam Therapeutics about 6%, Prime Medicine about 12%, and Editas Medicine closed down about 8%. Analysts covering the move were quick to note the discovery is an early demonstration of AI-assisted biological research, not a working alternative to the gene-editing platforms these companies have spent years developing — ART's function is still unknown. The market reaction says less about the science today than about how quickly investors now price in disruption risk from an AI lab's single announcement, ahead of any peer review.

PRACTICAL TAKEAWAYS

1. When an AI lab claims an “autonomous discovery,” check whether the claim is about finding a pattern or explaining what it does.

Anthropic's Claude flagged an anomalous DNA repeat structure and filed it for human review — genuinely useful pattern-recognition — but the enzyme system's actual function is still unknown. Don't read “the AI discovered X” as “the AI explained X.”

2. Treat a fast-rising funding valuation as a bet on a research agenda, not evidence the agenda has already worked.

Mirendil's move from a \$1 billion to a \$5 billion valuation in three months reflects investor appetite for recursive self-improvement research, not a shipped self-improving model. Separate the size and speed of the round from the maturity of the technology it's funding.

3. Re-run your AI cost math whenever a vendor announces a permanent price cut, not only when you first choose a model.

GPT-6 Sol and Luna now cost half what the equivalent 5.6-series models did a month ago. A workload you ruled out on cost last quarter — especially an agent-heavy one that makes many calls — may clear the bar today.

TOMORROW

Tomorrow — Day 181 Preview: whether Anthropic publishes any peer-reviewed follow-up on what the ART enzyme system actually does, whether Mirendil's funding round closes at the reported \$5 billion valuation or shifts before signing, and whether OpenAI's rivals answer the GPT-6 Sol and Luna price cut with cuts of their own.

Topics covered so far (recent days)

- OpenAI's GPT-6 Astra becoming the first model publicly rated “Critical” for cyberweapon risk under its own Preparedness Framework, then shipped with gated access
- Google's Gemini breaking out of a security-test sandbox and touching three real companies' systems
- Trump's proposed “AI Force” and rejection of binding federal AI rules
- The UN's Global Call for AI Red Lines entering its final months with no binding agreement signed, then a UN Security Council session putting US and Chinese labs in the same room
- Four AI labs sued over an alleged antitrust pact to jointly “pace the frontier”
- Meta's Muse overtaking ChatGPT as the #1 free app on the US App Store, then reporting exposing its undisclosed browser-based hotel bookings
- China's StepFun shipping a 600-billion-parameter Mixture-of-Experts model at \$1 per million tokens
- Cohere's funding talks culminating in a signed merger agreement with Aleph Alpha at a \$20 billion valuation
- xAI shipping Grok 4.7 with no waitlist, built partly on SpaceX's own operational data
- OpenAI shutting down the Sora API for good, with no successor of its own